

# OPENSIFT CONTAINER PLATFORM

## ARCHITECTURAL OVERVIEW



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[facebook.com/redhatinc](https://www.facebook.com/redhatinc)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



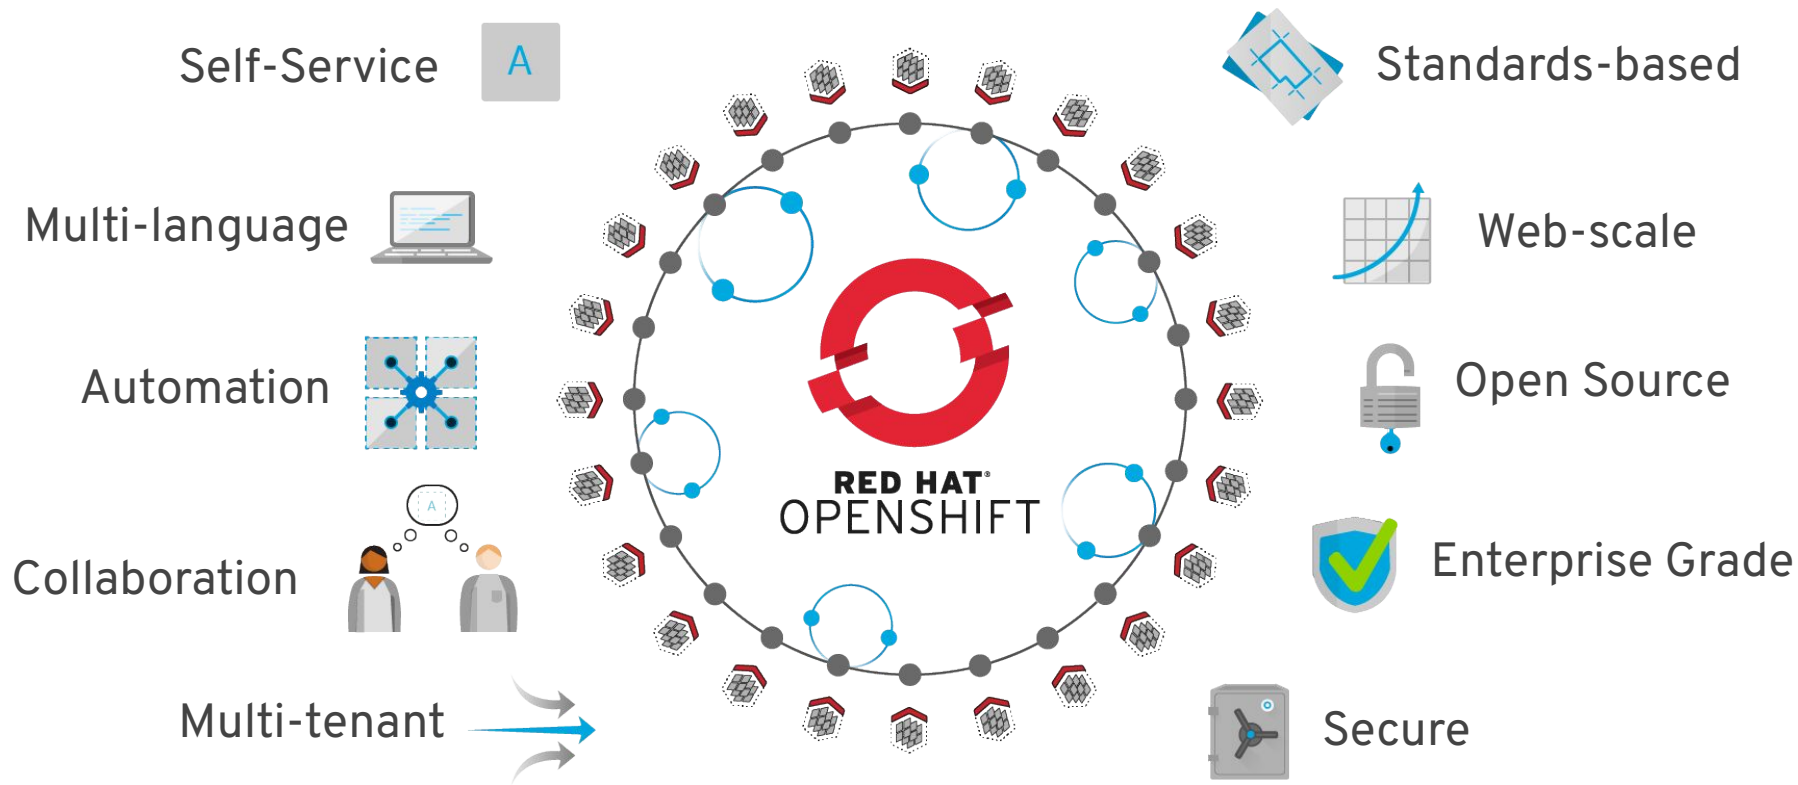
[twitter.com/RedHat](https://twitter.com/RedHat)

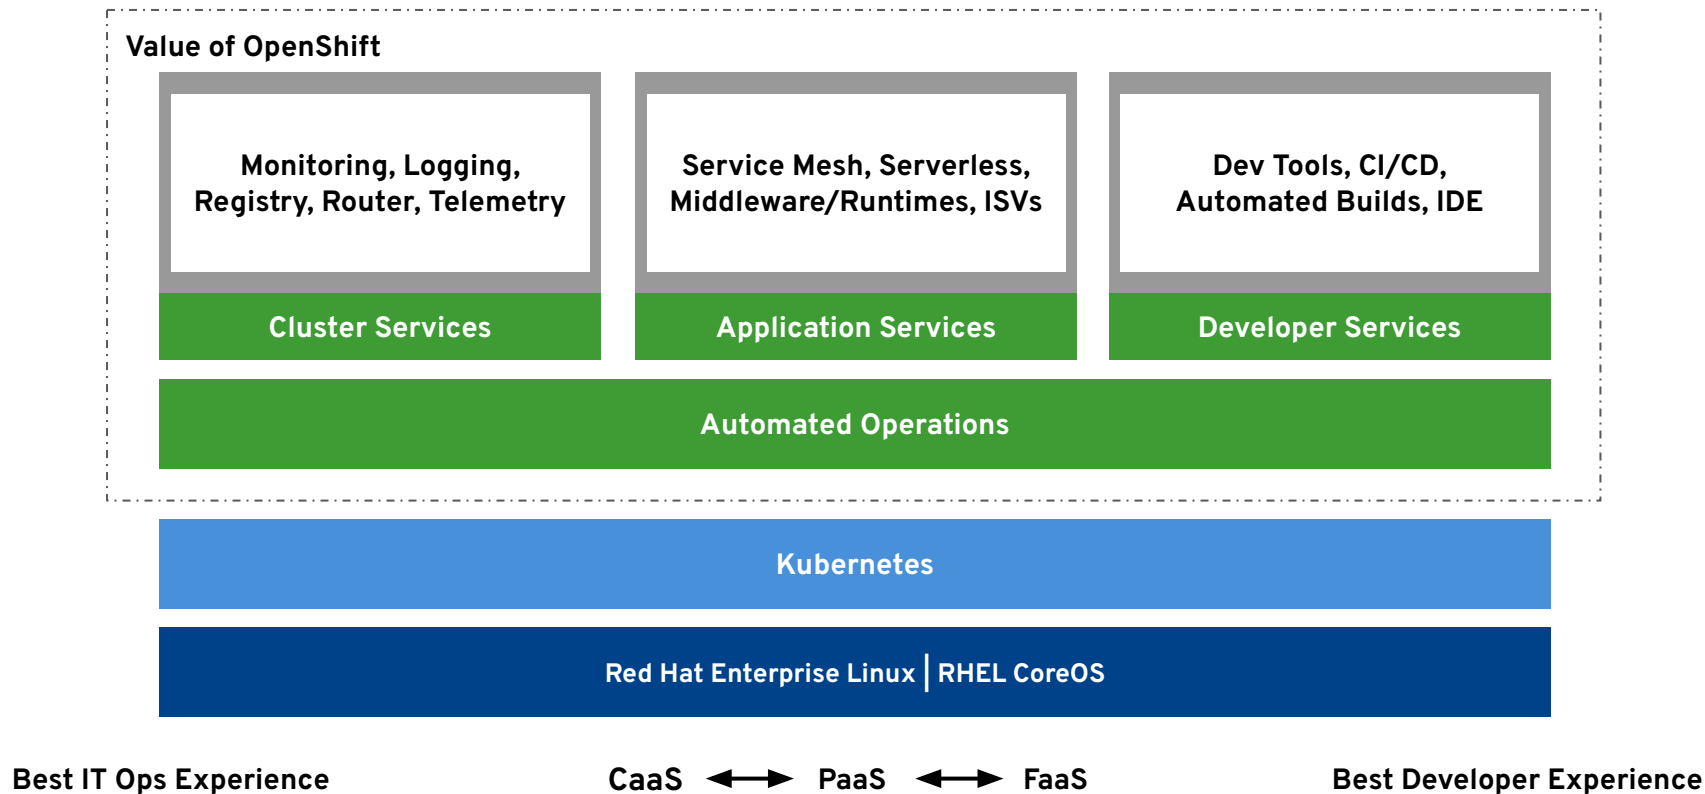
Alfred Bach  
Principal Solution Architect  
Feb 2022

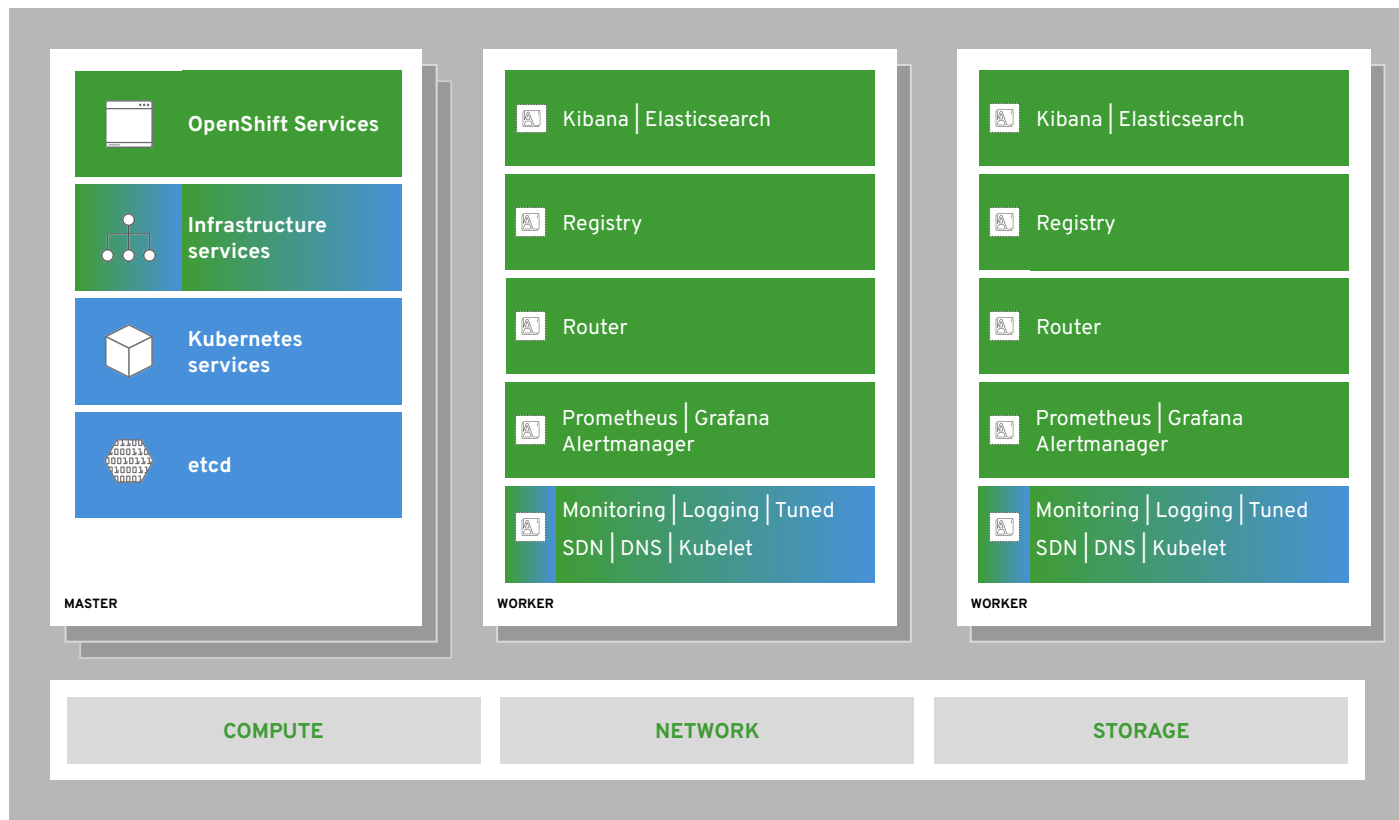
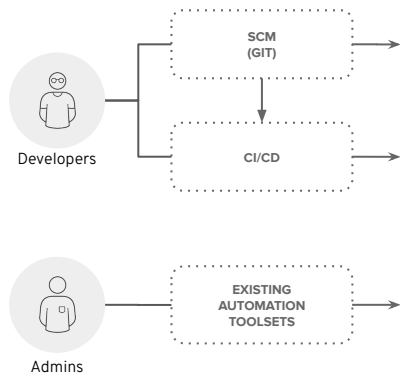




# Functional overview







Database

Streaming & Messaging

Application Definition & Image Build

Continuous Integration & Delivery

Platform

Observability and Analysis

App Definition and Development



Orchestration & Management



Cloud-Native Storage

Container Runtime

Cloud-Native Network

Runtime



Automation & Configuration

Container Registry

Security & Compliance

Key Management

Provisioning



Public

Kubernetes Certified Service Provider

Kubernetes Training Partner


Cloud Native Computing Foundation

Cloud Native Landscape

l.cncf.io

This landscape is intended as a map through the previously uncharted terrain of cloud native technologies. There are many routes to deploying a cloud native application, with CNCF Projects representing a particularly well-traveled path.





# OpenShift and Kubernetes core concepts

# a container is the smallest compute unit

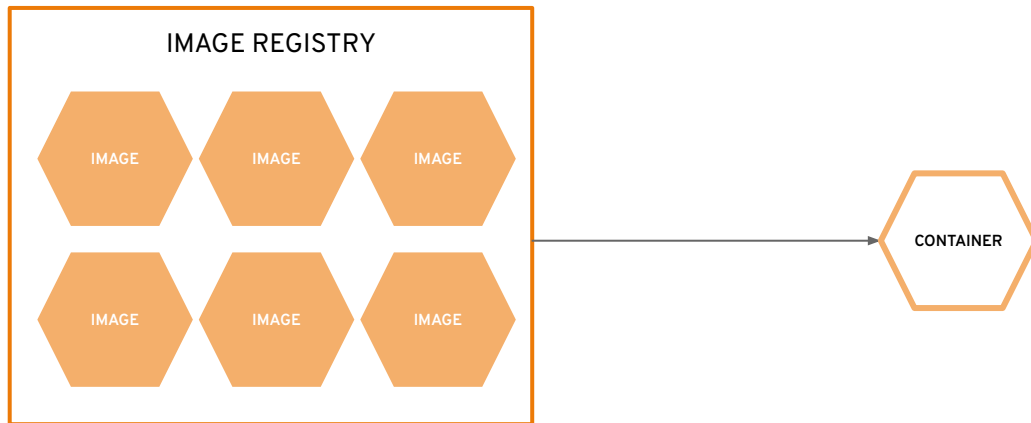




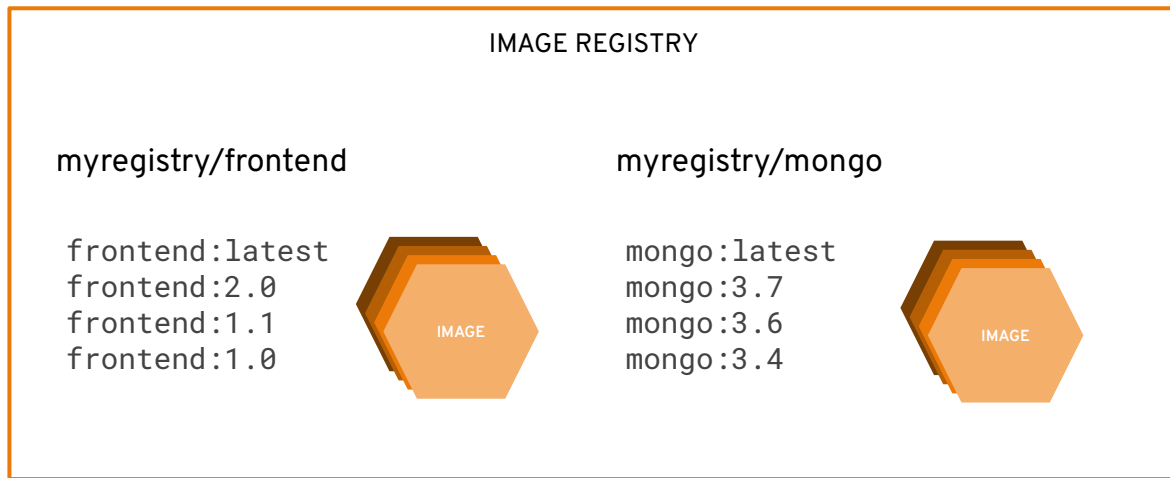
# containers are created from container images



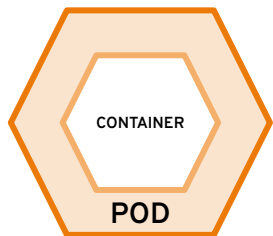
# container images are stored in an image registry



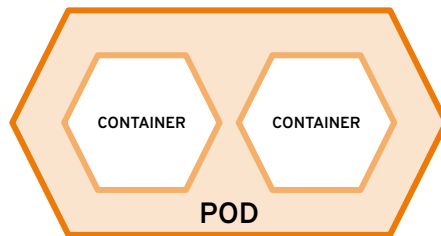
# an image repository contains all versions of an image in the image registry



# containers are wrapped in pods which are units of deployment and management

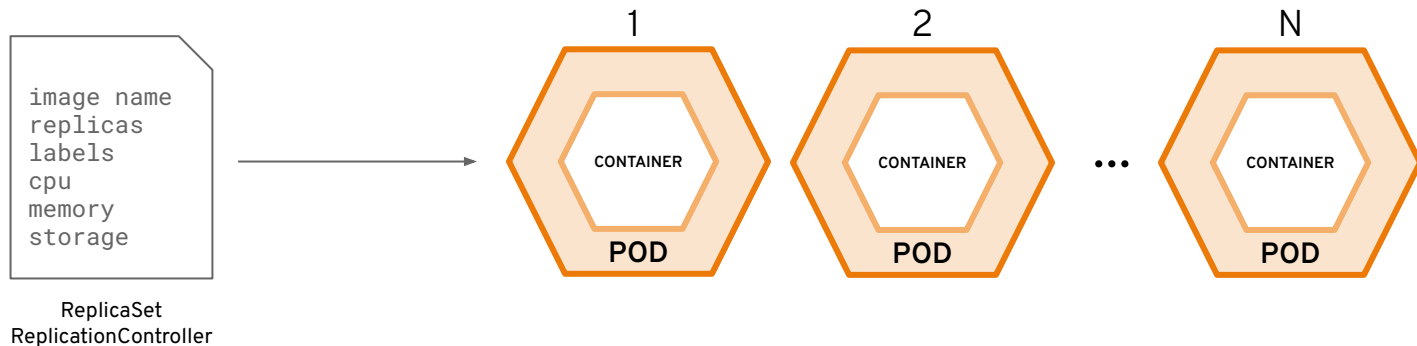


10.140.4.44

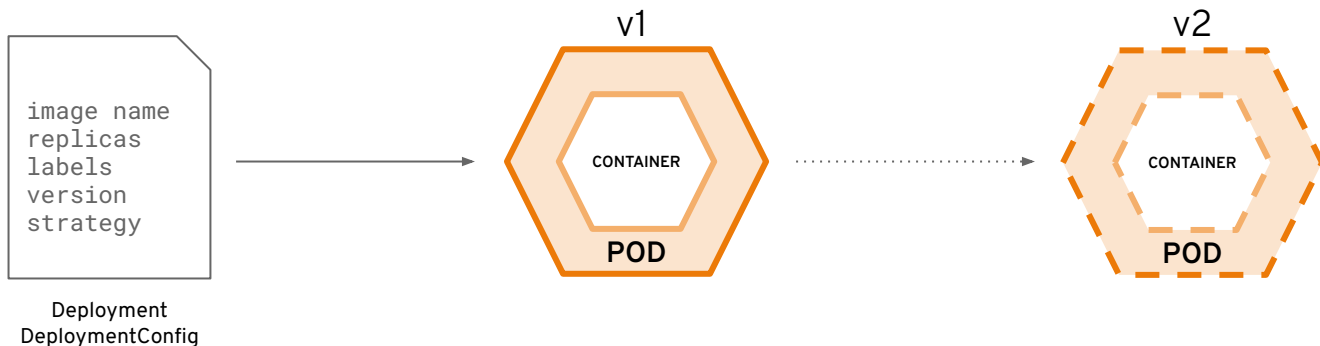


10.15.6.55

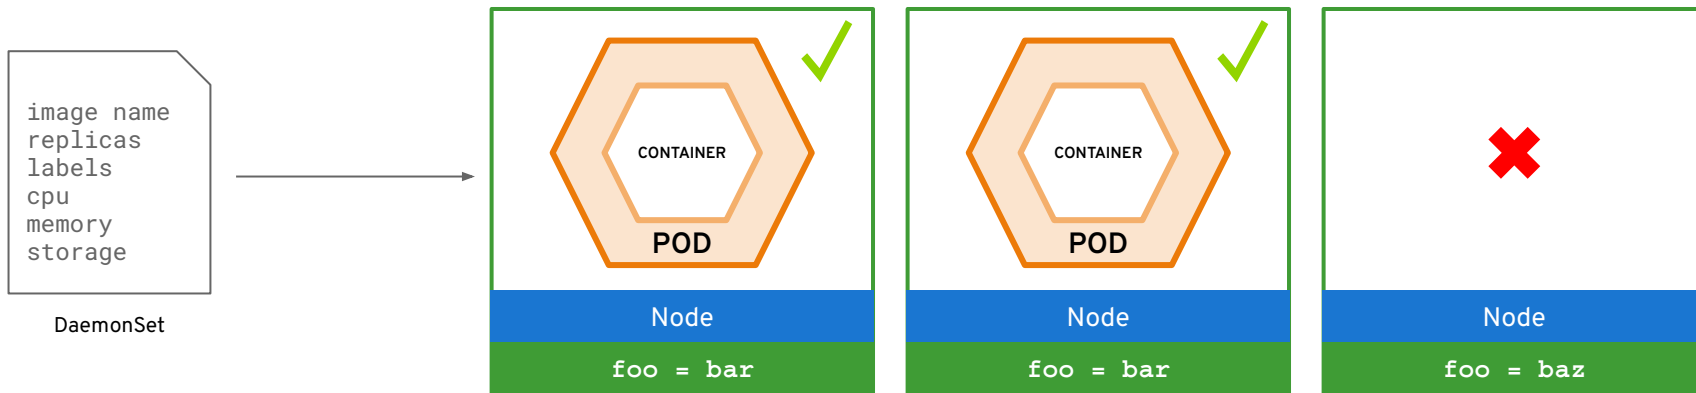
# ReplicationControllers & ReplicaSets ensure a specified number of pods are running at any given time



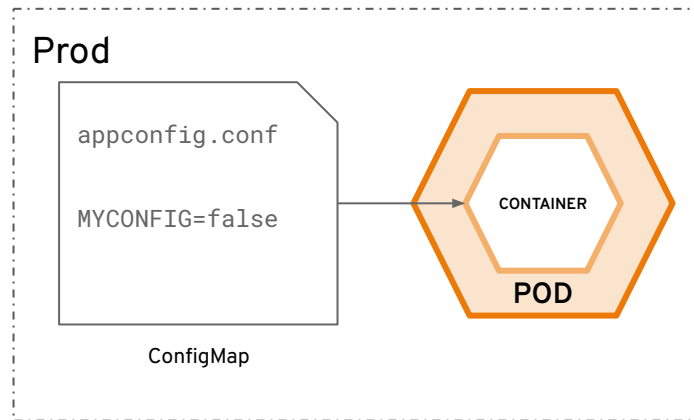
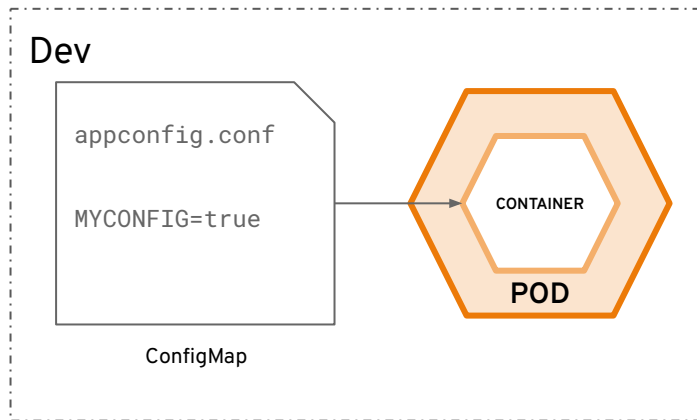
# Deployments and DeploymentConfigurations define how to roll out new versions of Pods



a daemonset ensures that all  
(or some) nodes run a copy of a  
pod

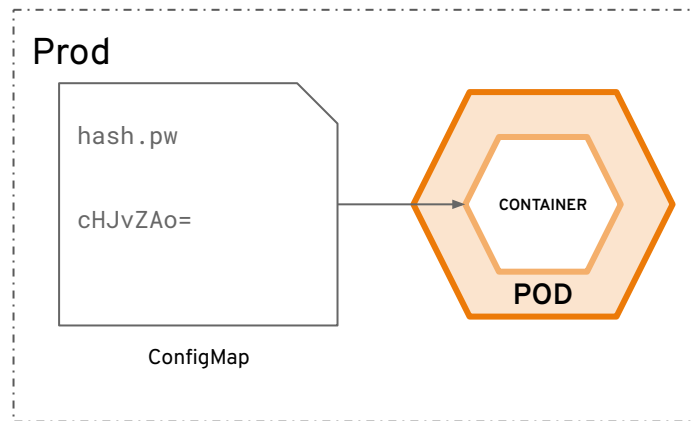
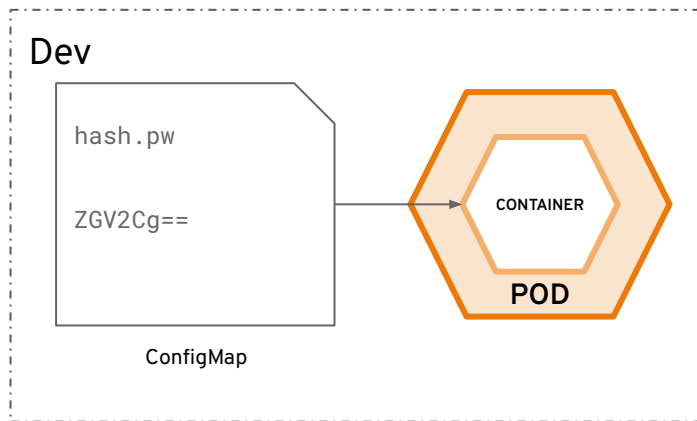


# configmaps allow you to decouple configuration artifacts from image content





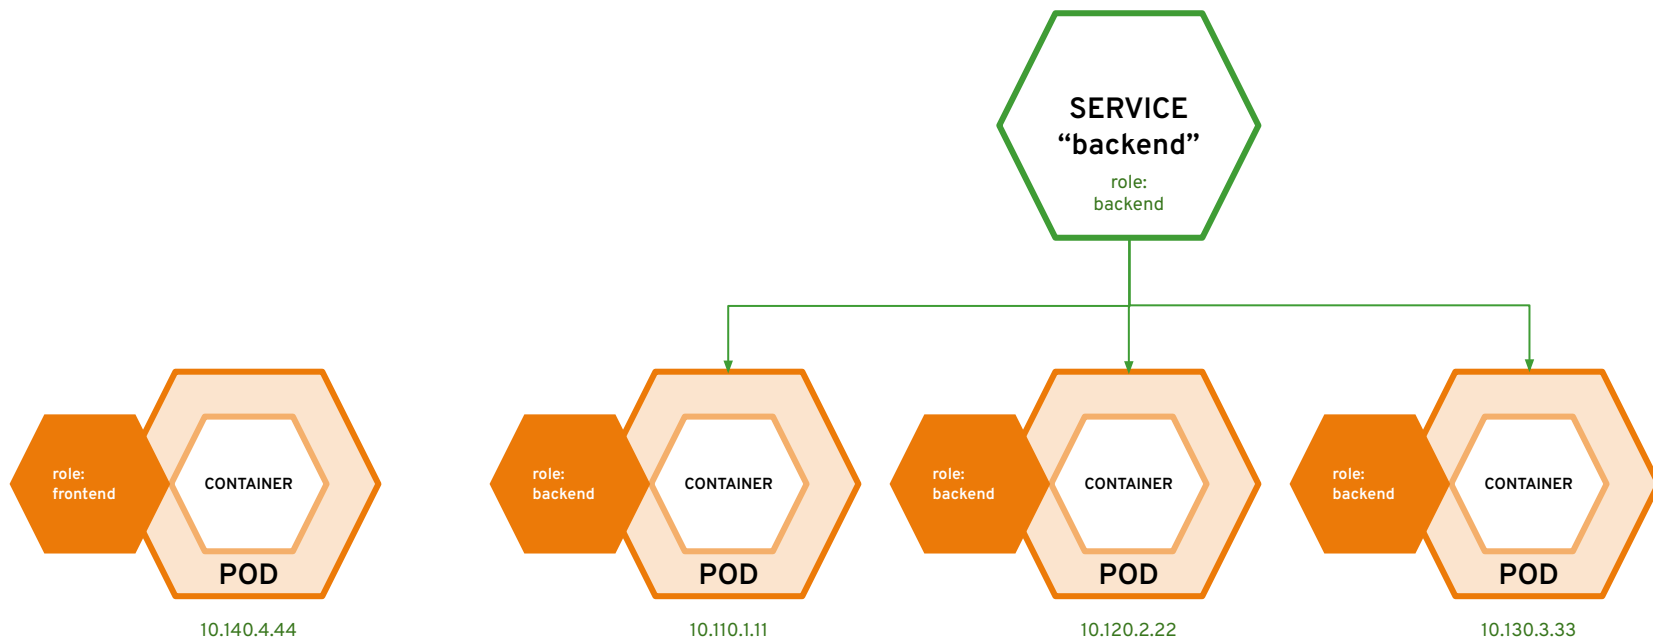
secrets provide a mechanism to hold sensitive information such as passwords



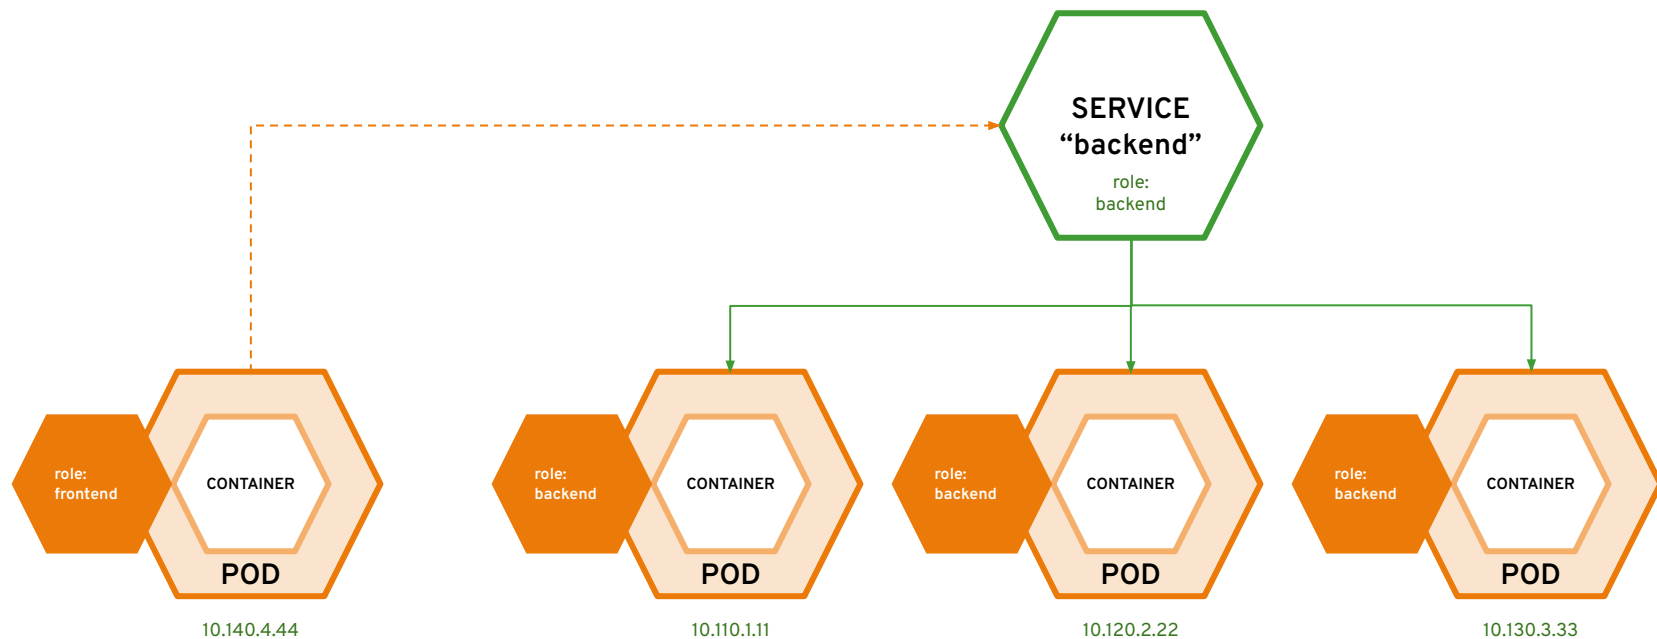
The etcd datastore can be encrypted for additional security

<https://docs.openshift.com/container-platform/4.6/security/encrypting-etcd.html>

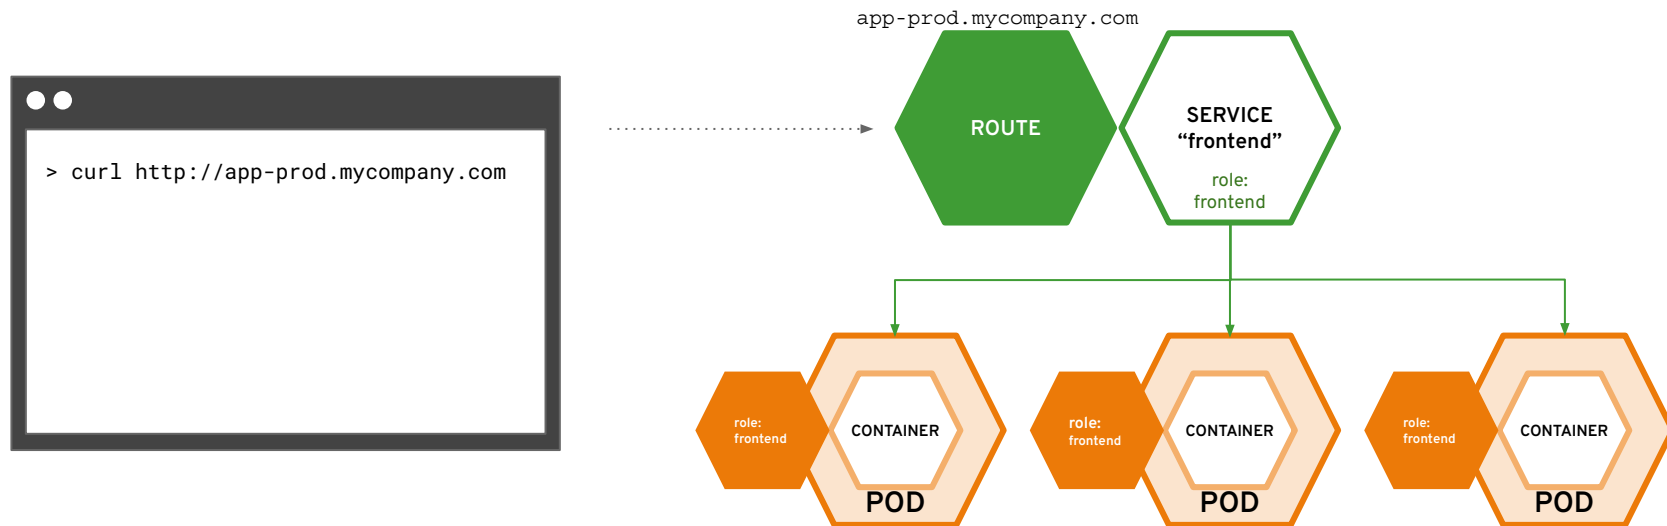
# services provide internal load-balancing and service discovery across pods



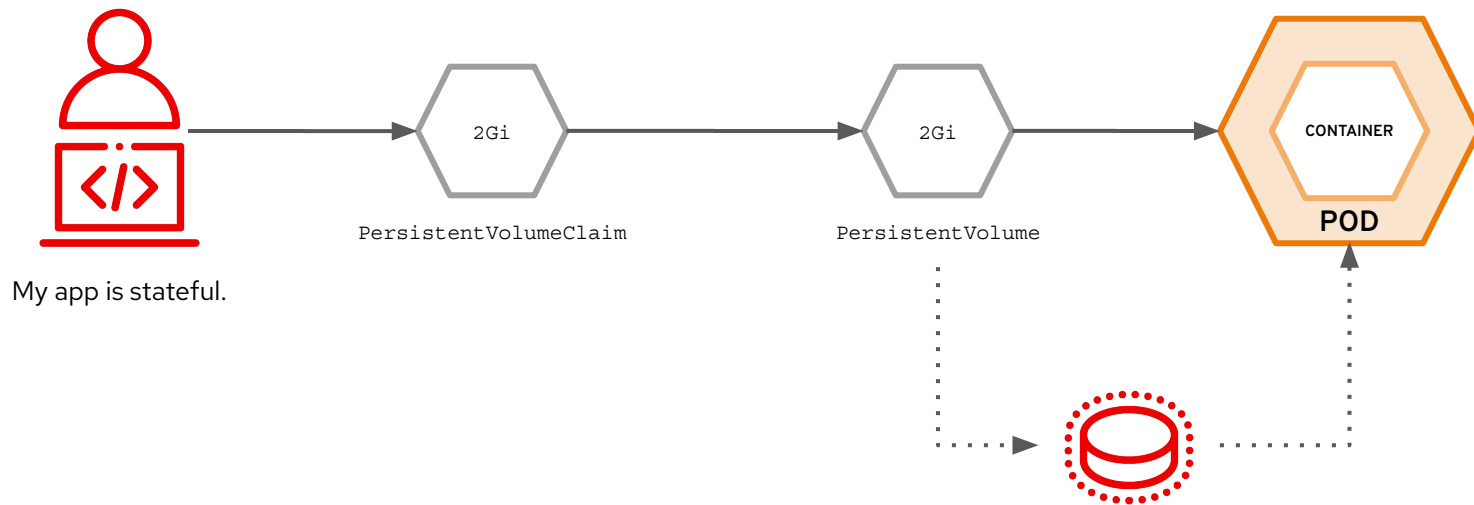
# apps can talk to each other via services



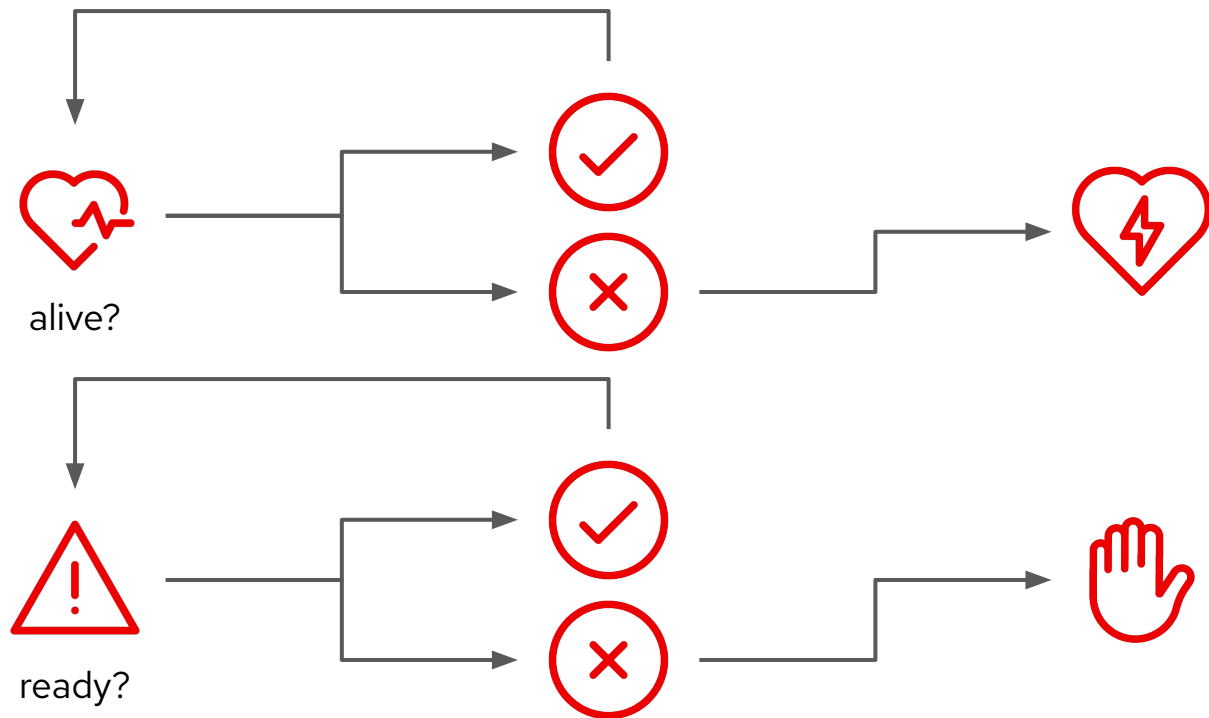
routes make services accessible to clients outside the environment via real-world urls



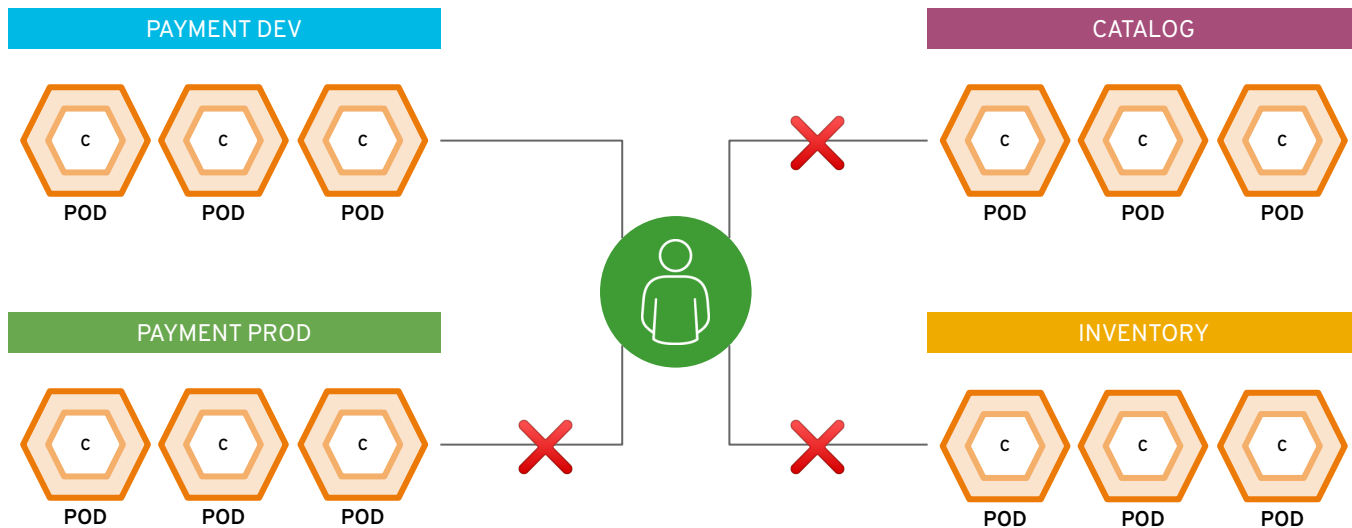
# Persistent Volume and Claims



# Liveness and Readiness



# projects isolate apps across environments, teams, groups and departments





# OpenShift 4 Architecture

We start again at 10:15 CEST



# your choice of infrastructure

COMPUTE

NETWORK

STORAGE

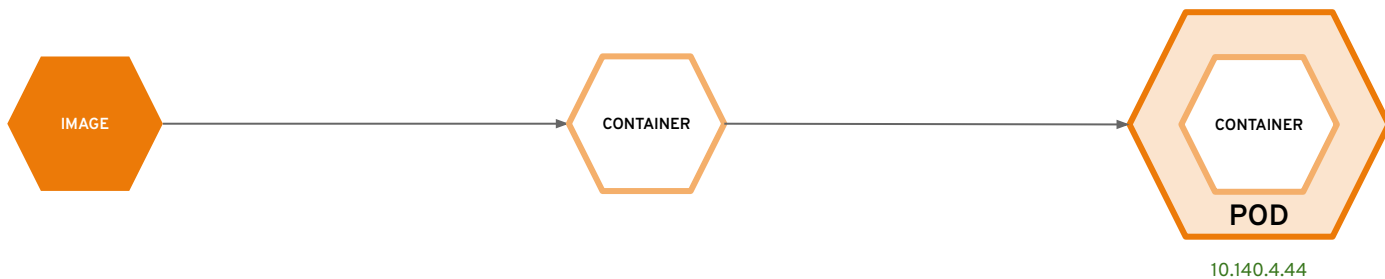
## workers run workloads



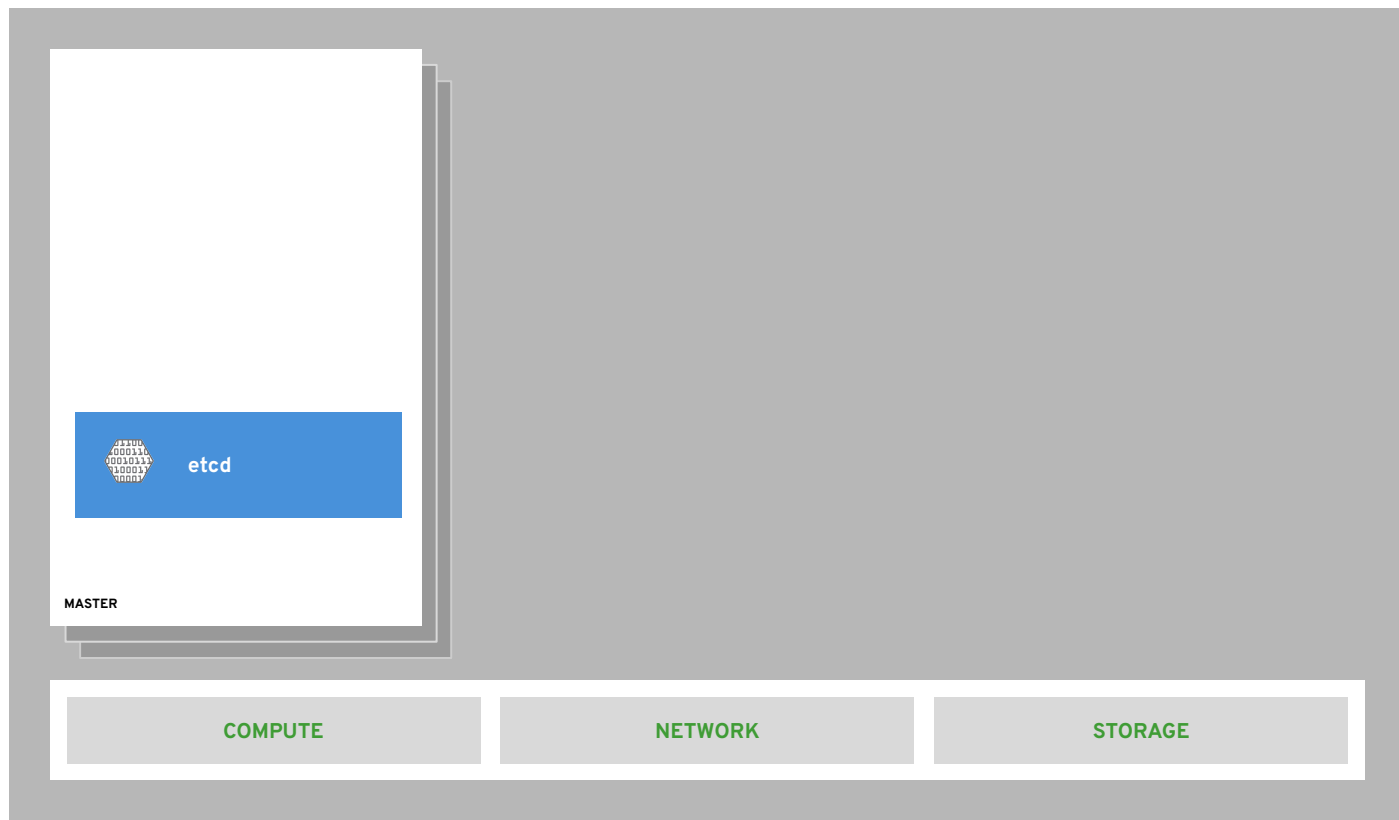
# masters are the control plane



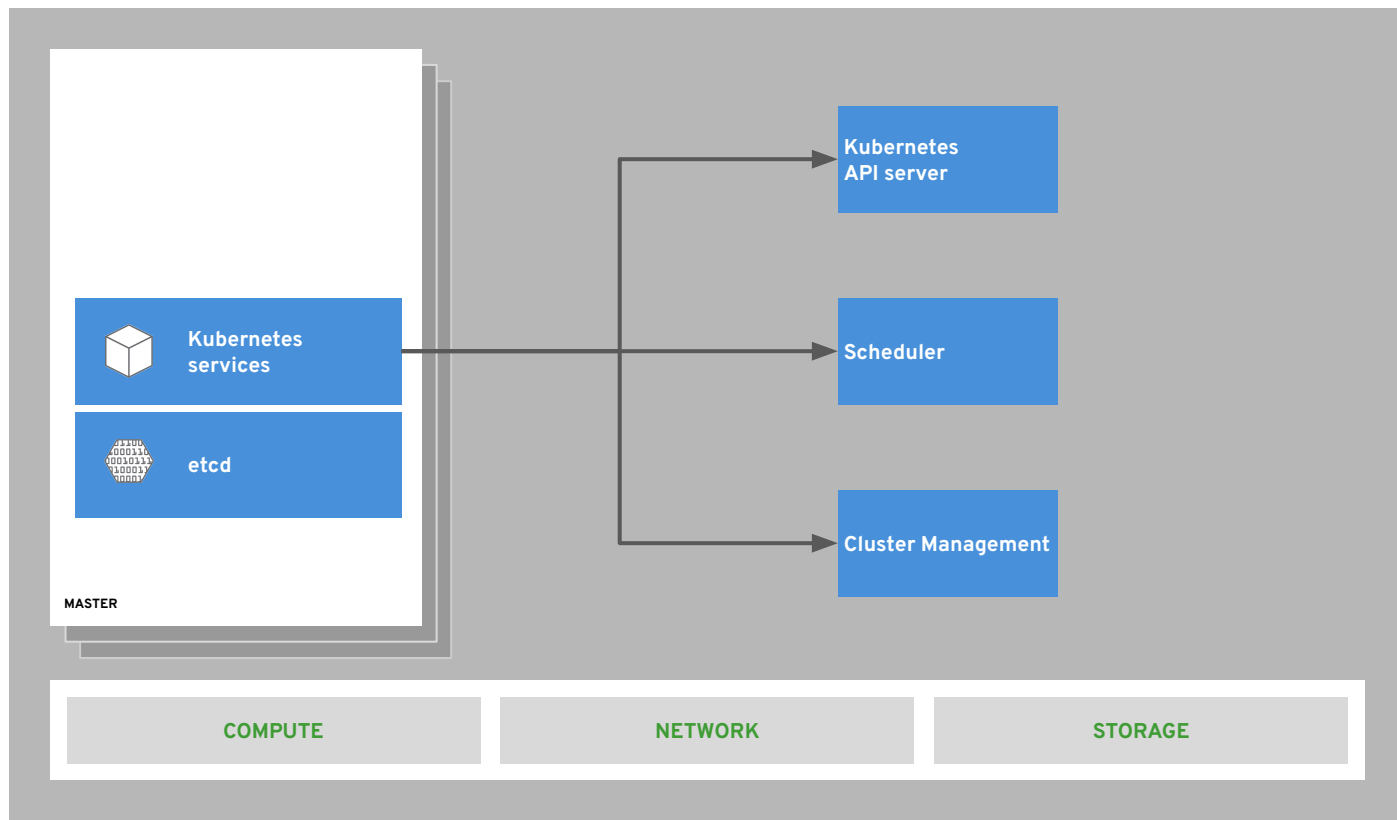
# everything runs in pods



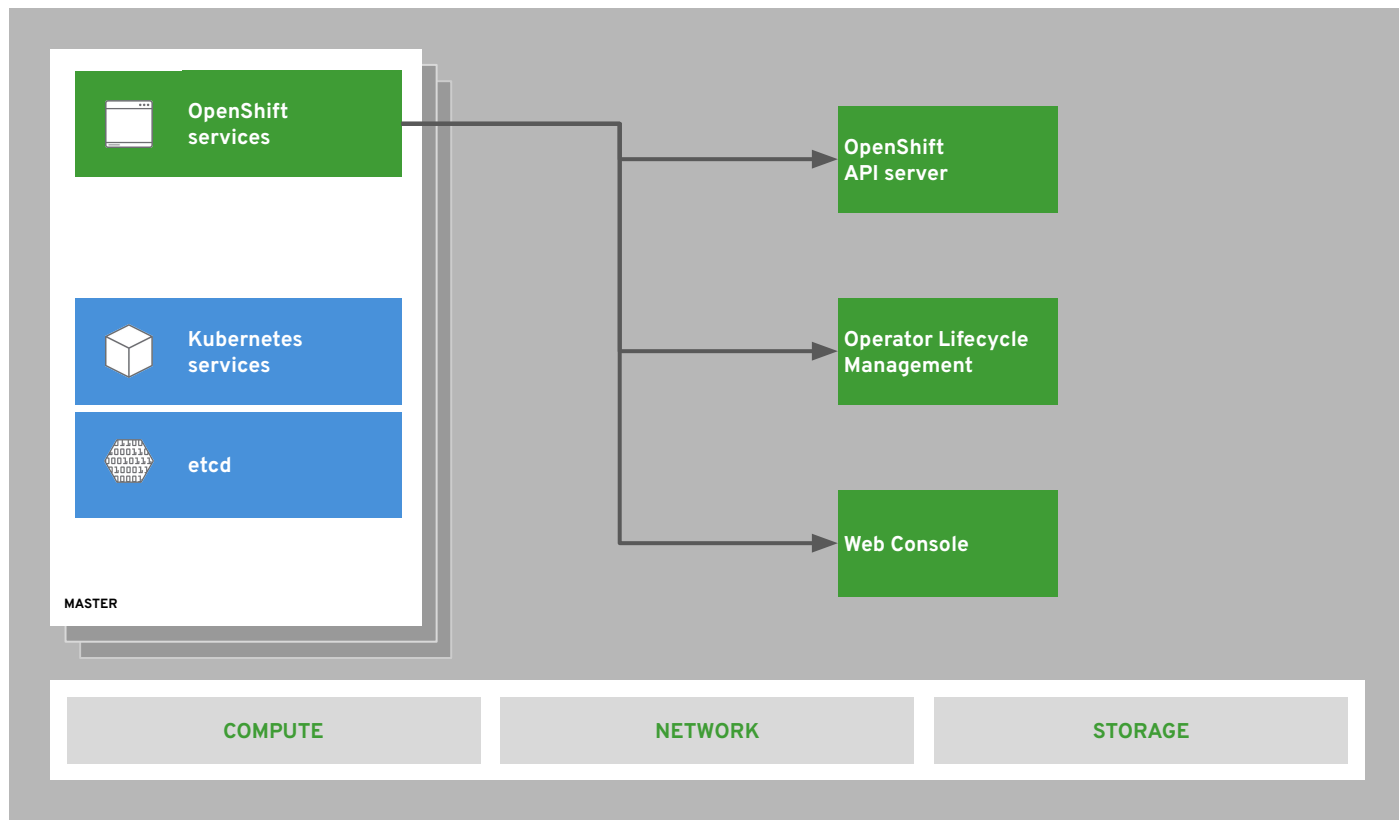
# state of everything



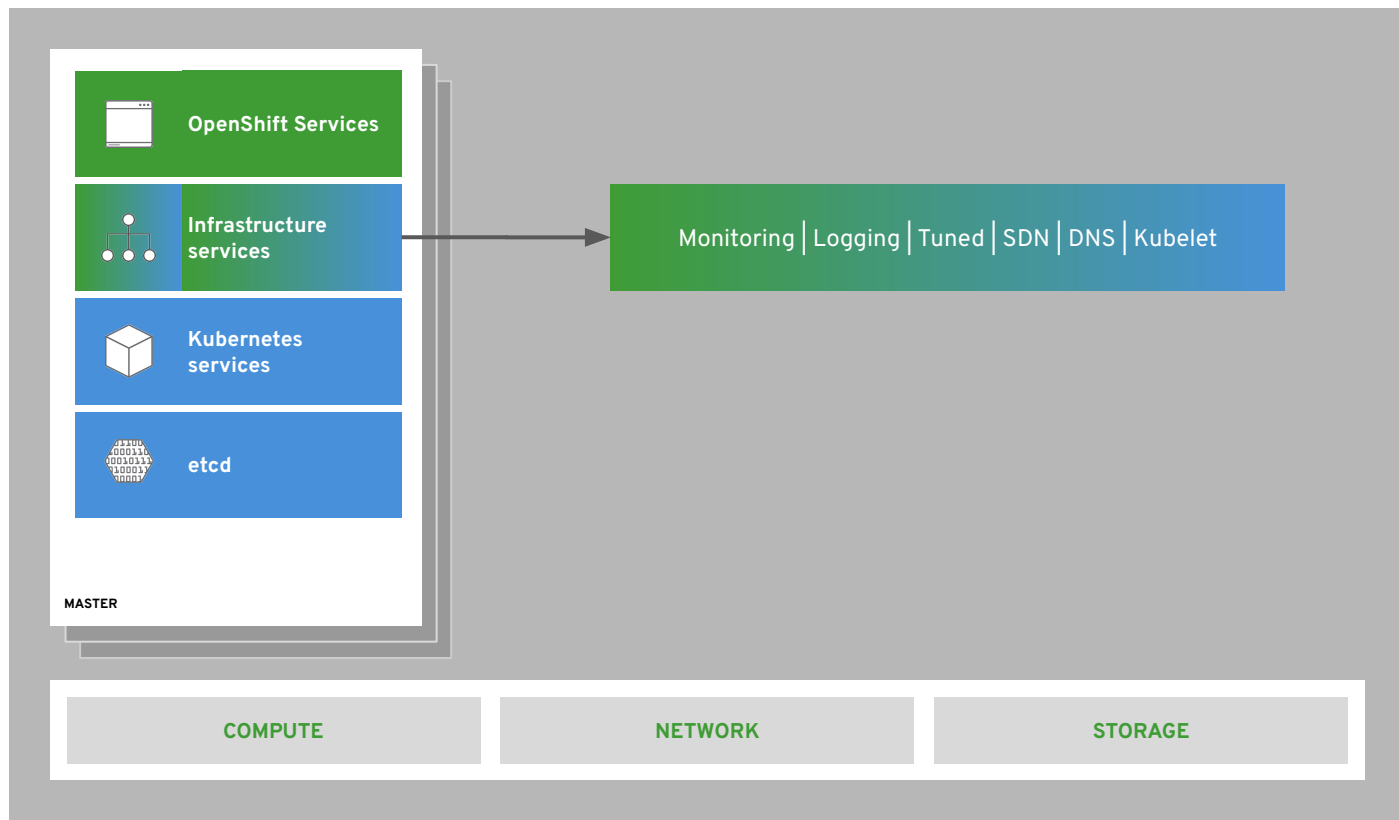
# core kubernetes components



# core OpenShift components

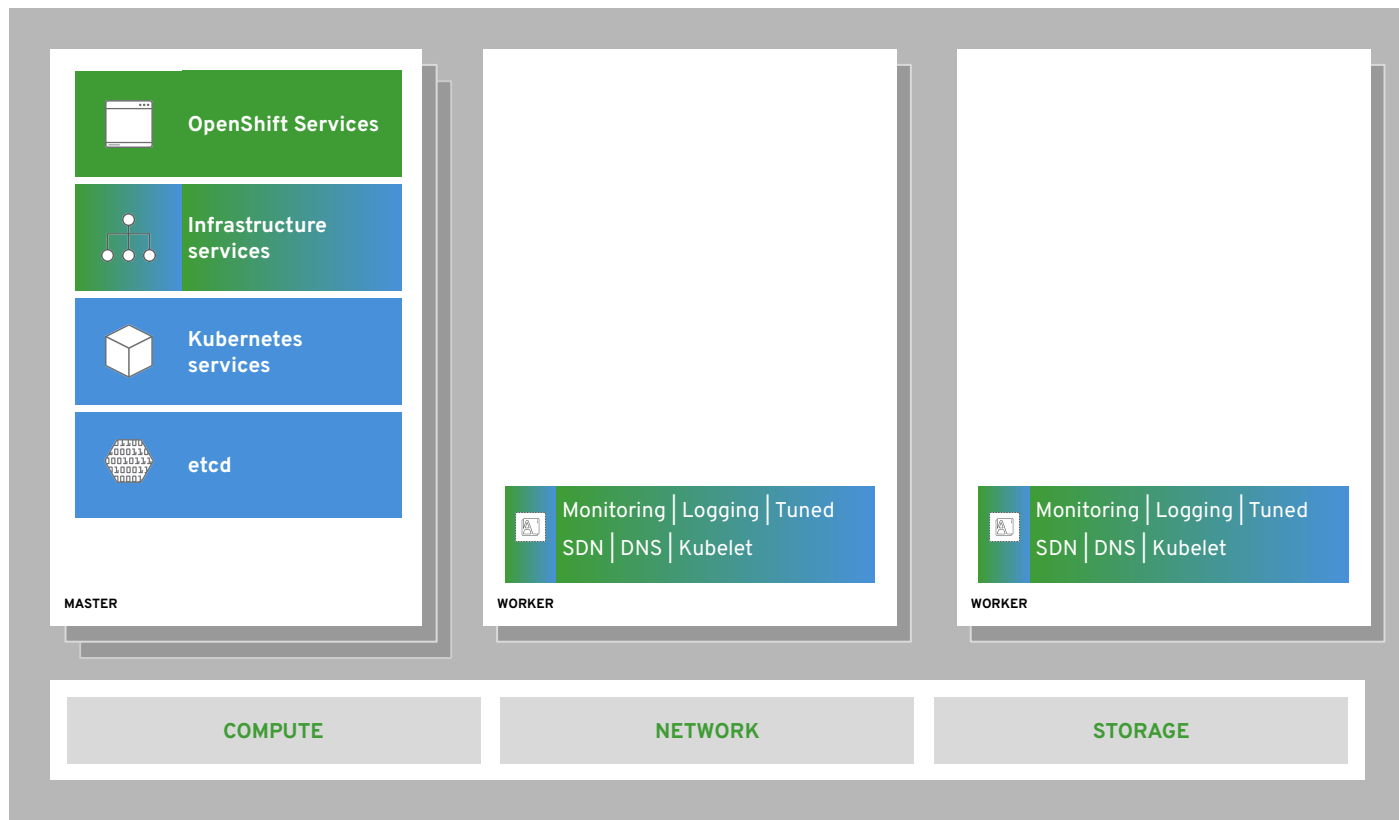


# internal and support infrastructure services

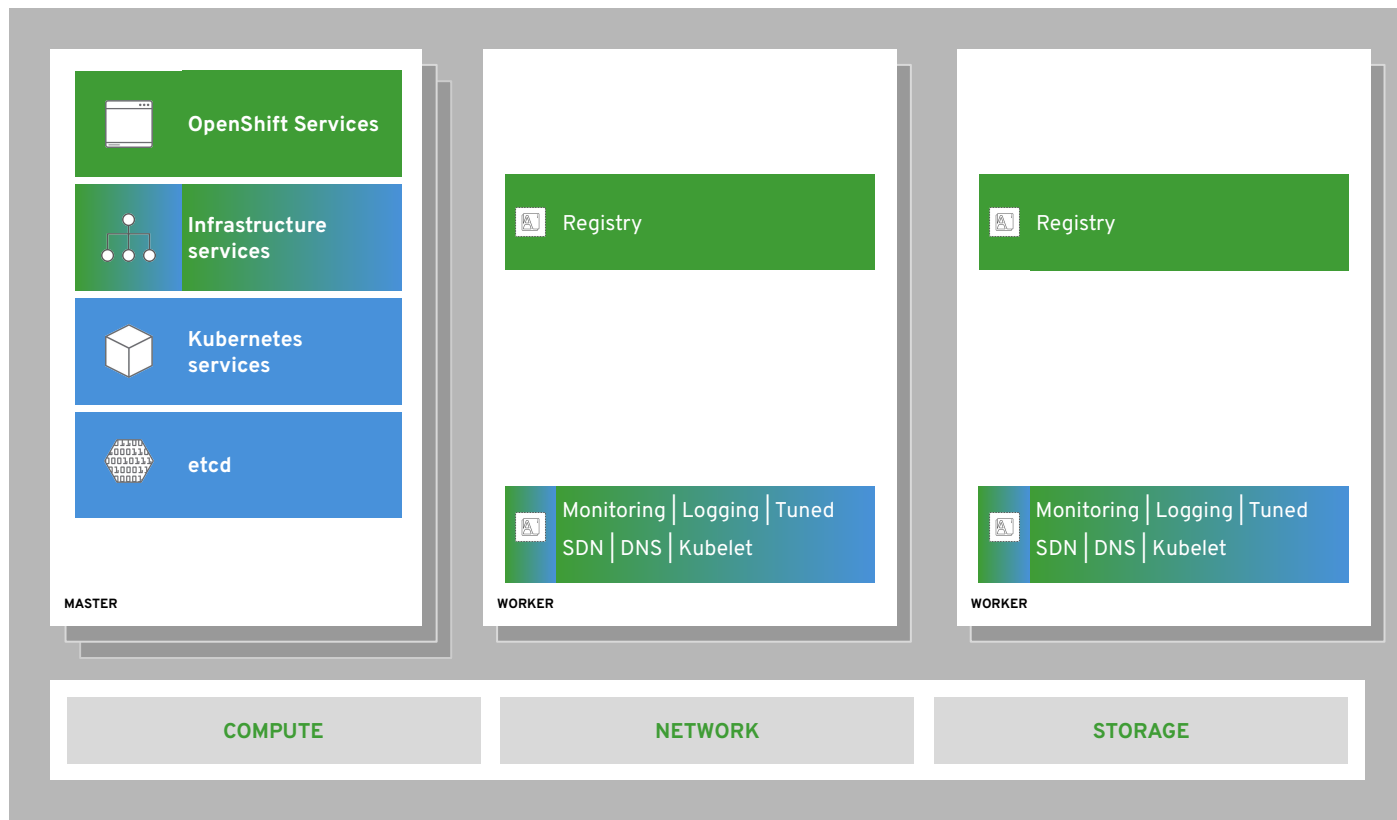




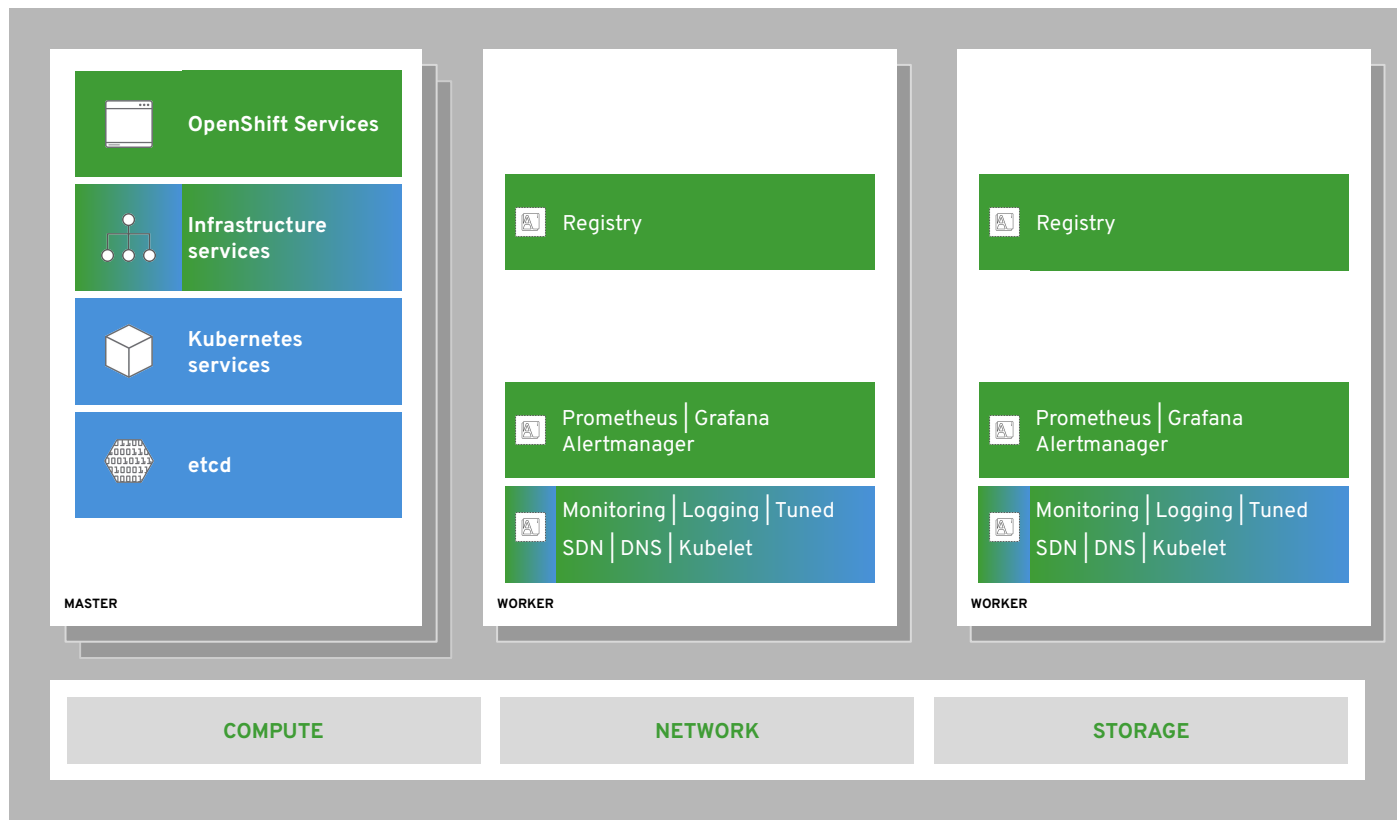
# run on all hosts



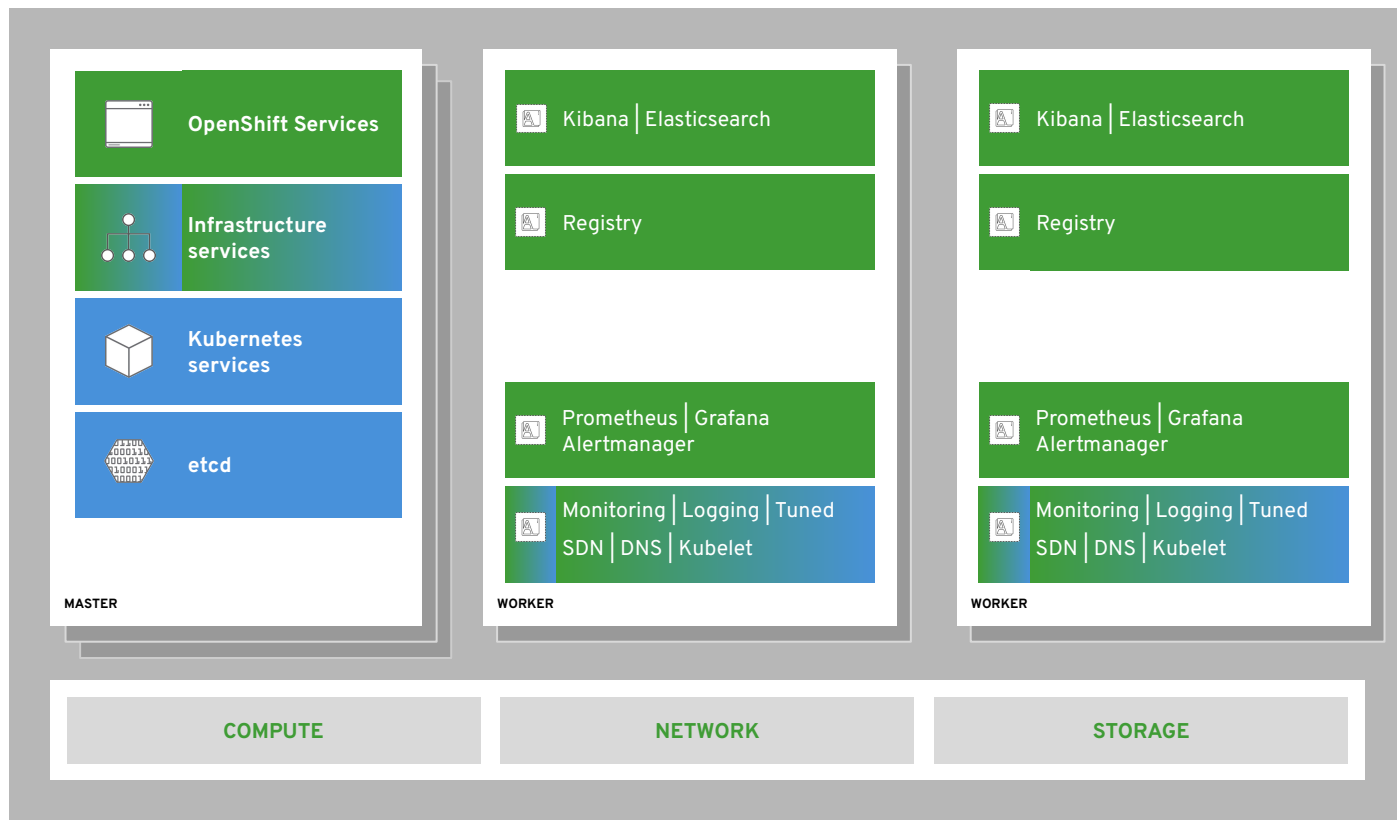
# integrated image registry



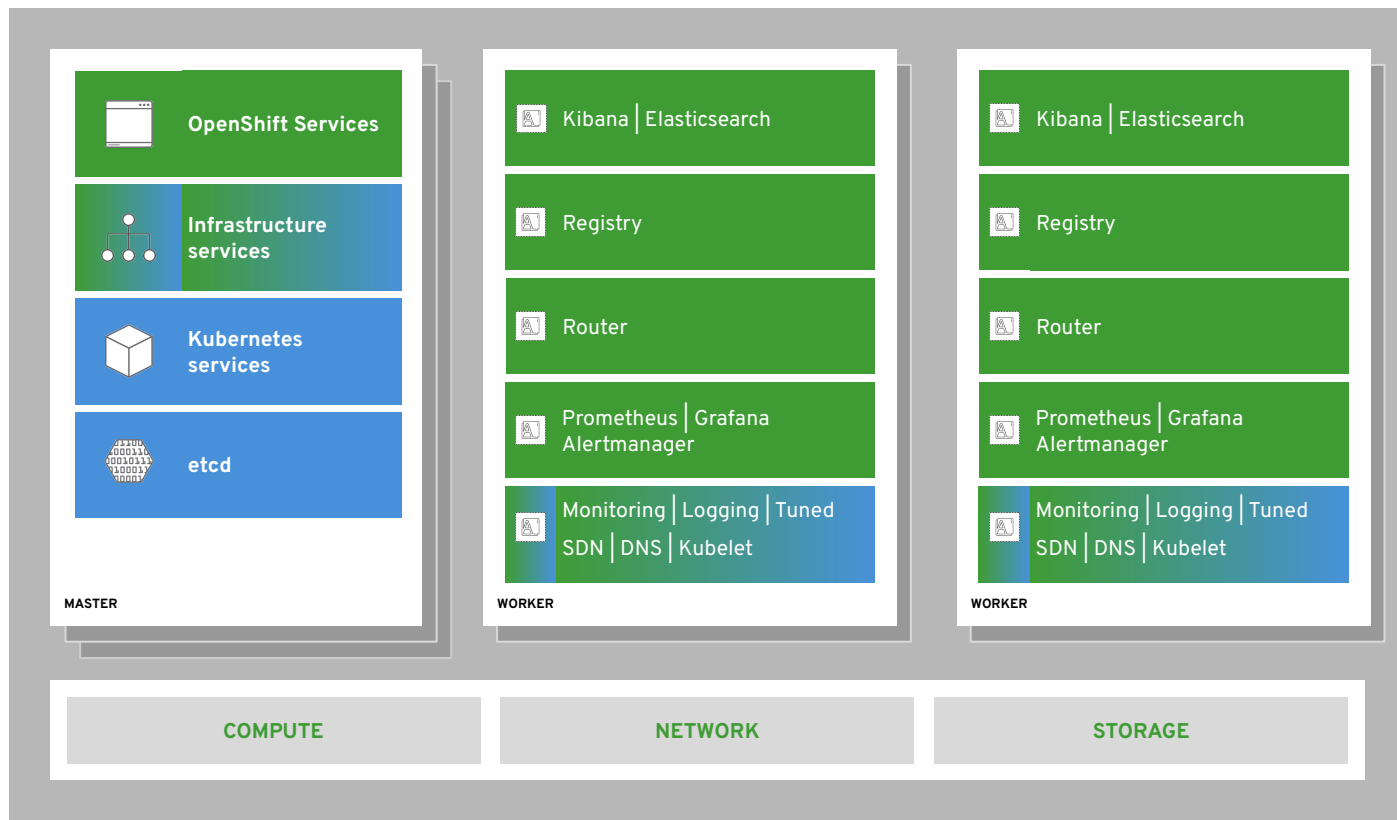
# cluster monitoring



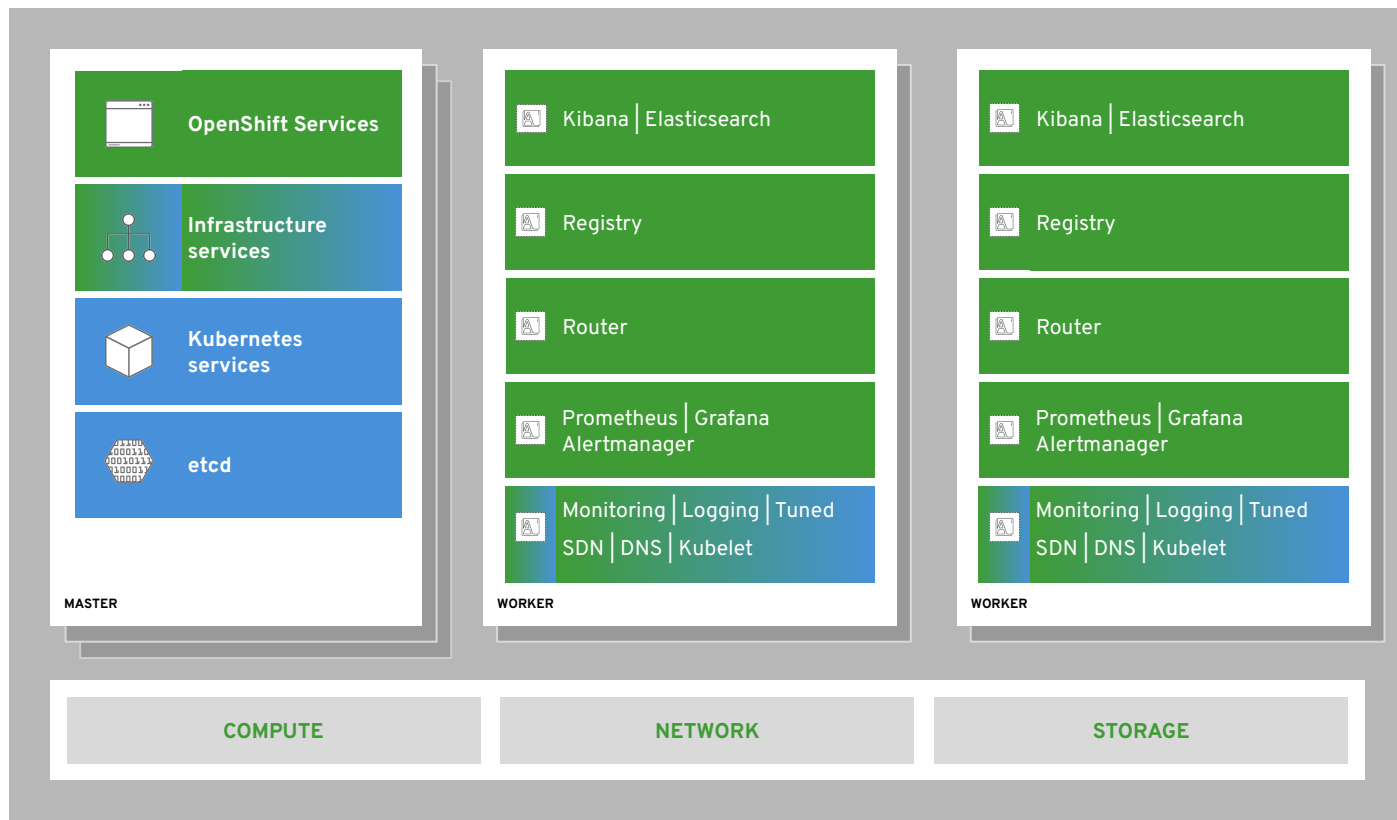
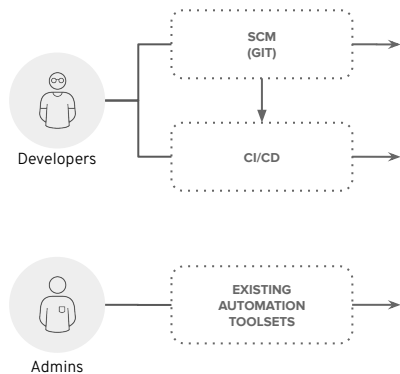
# log aggregation



# integrated routing



# dev and ops via web, cli, API, and IDE



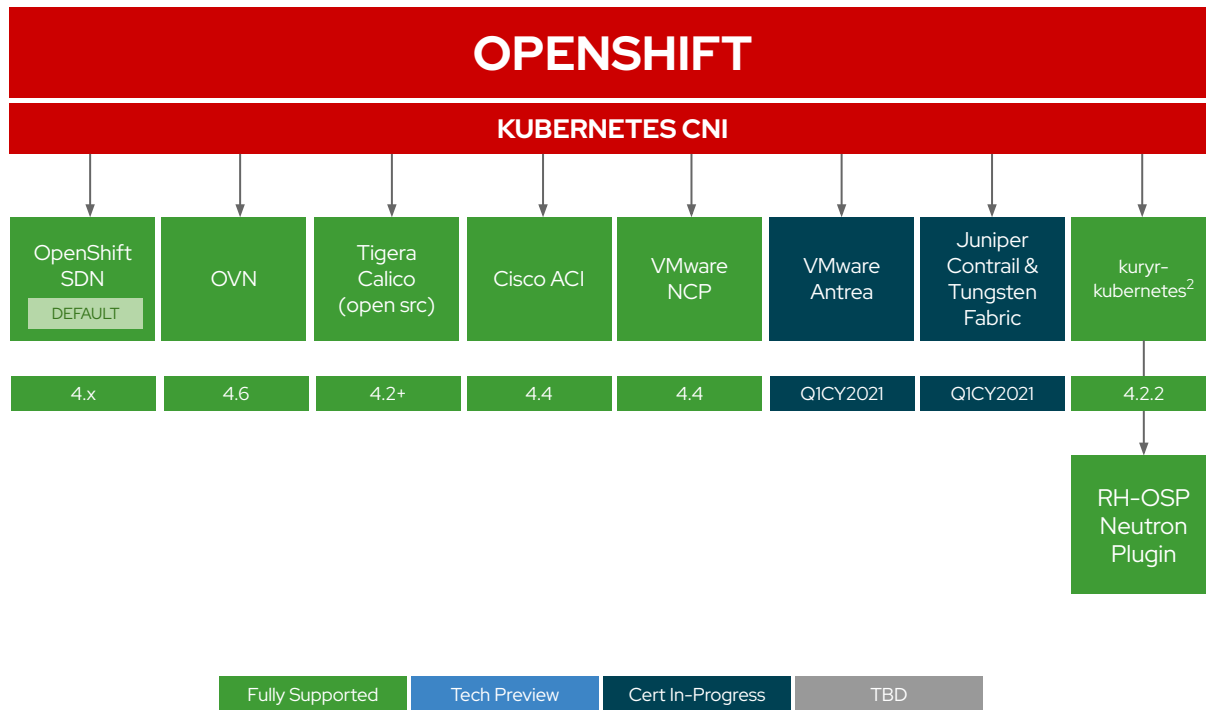
# Networking

A pluggable model for  
network interface  
controls in kubernetes

# OpenShift Networking Plug-ins

3rd-party Kubernetes CNI plug-in certification primarily consists of:

1. Formalizing the partnership
2. Certifying the container(s)
3. Certifying the Operator
4. Successfully passing the same Kubernetes networking conformance tests that OpenShift uses to validate its own SDN

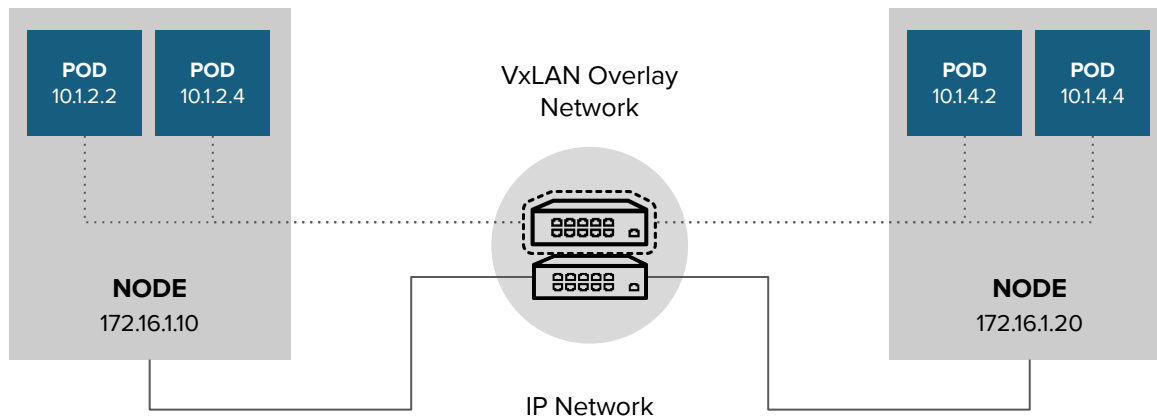




# OpenShift SDN

An Open  
vSwitch-based  
Software Defined  
Network for  
kubernetes

# OpenShift SDN high-level architecture



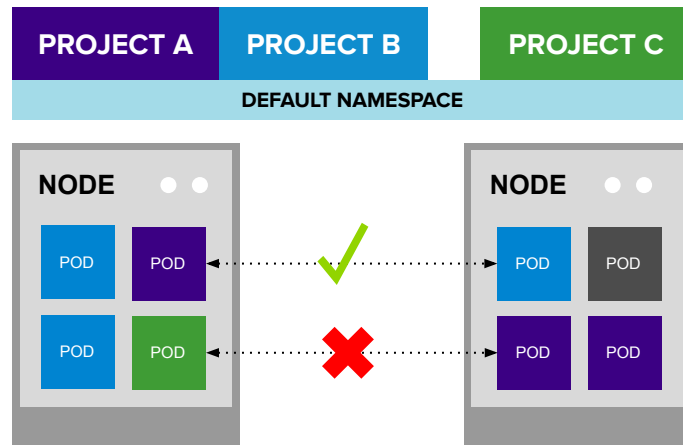
# OpenShift SDN “flavors”

## OPEN NETWORK (Default)

- All pods can communicate with each other across projects

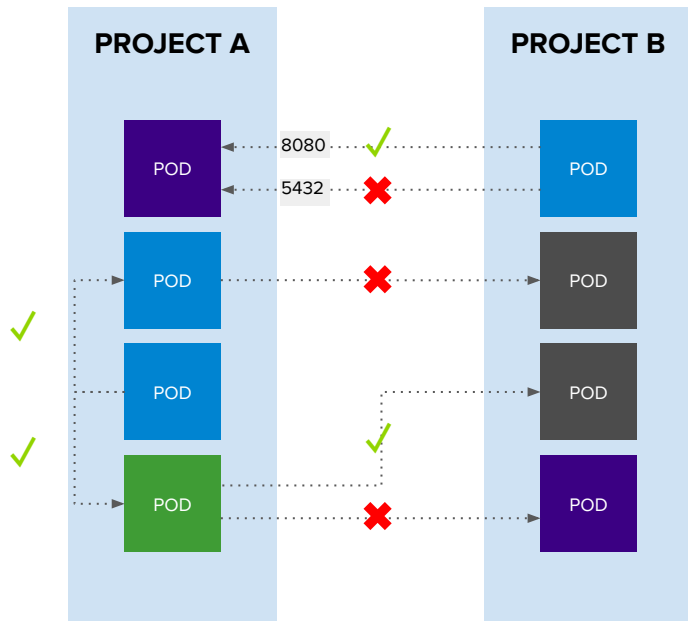
## MULTI-TENANT NETWORK

- Project-level network isolation
- Multicast support
- Egress network policies



Multi-Tenant Network

# NetworkPolicy

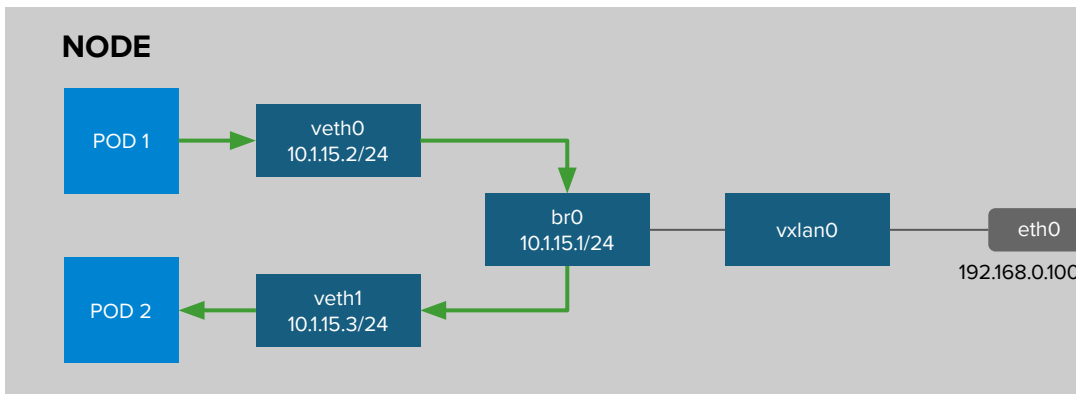


## Example Policies

- Allow all traffic inside the project
- Allow traffic from green to gray
- Allow traffic to purple on 8080

```
apiVersion: extensions/v1beta1
kind: NetworkPolicy
metadata:
  name: allow-to-purple-on-8080
spec:
  podSelector:
    matchLabels:
      color: purple
  ingress:
    - ports:
        - protocol: tcp
          port: 8080
```

# OpenShift SDN packet flows container-container on same host

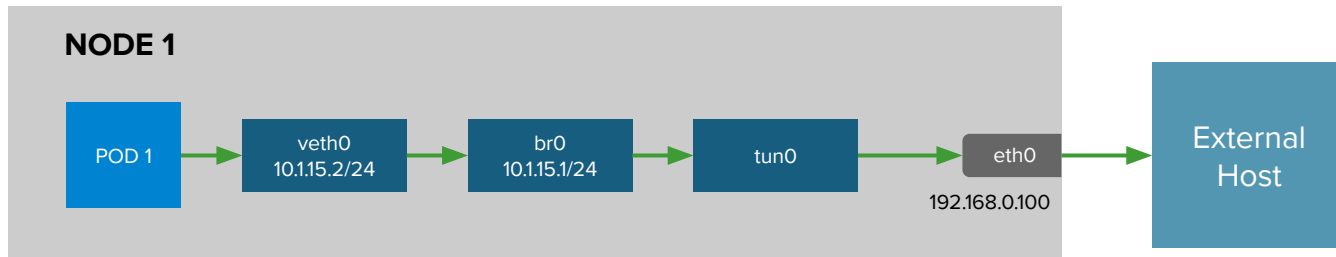


# OpenShift SDN packet flows

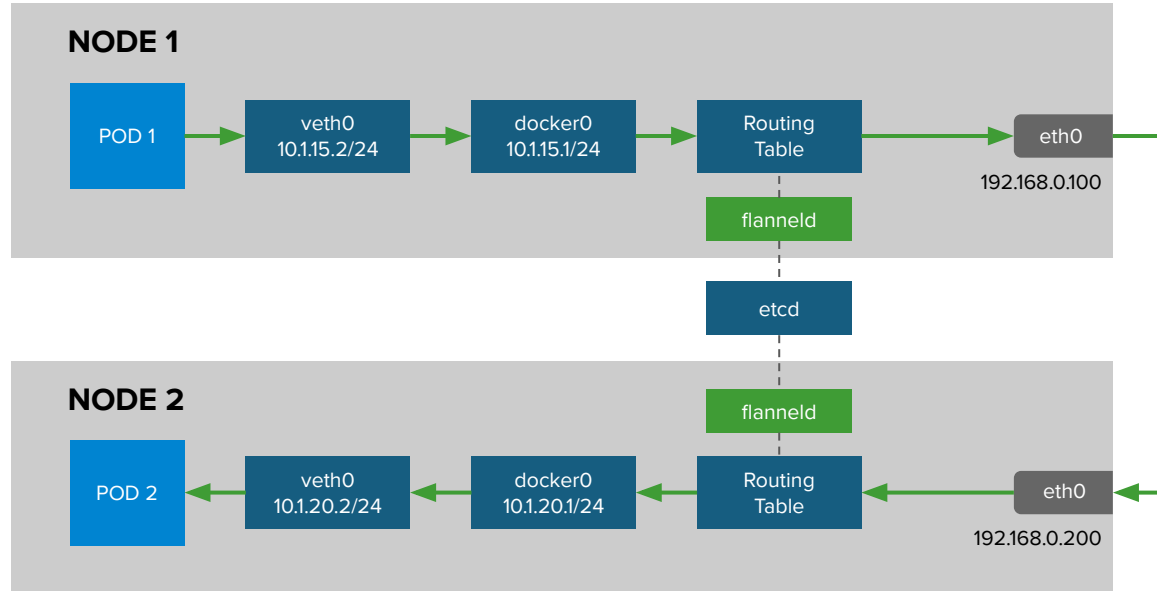
## container-container across hosts



# OpenShift SDN packet flows container leaving the host



# Kuryr and OpenStack



Flannel is minimally verified and is supported only and exactly as deployed in the OpenShift on OpenStack reference architecture <https://access.redhat.com/articles/2743631>

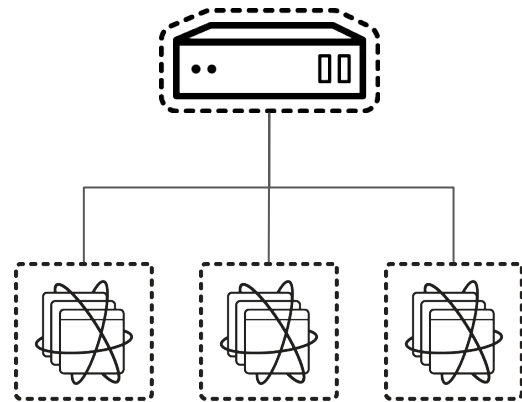


# routes and ingress

How traffic enters the  
cluster

## Routing and Load Balancing

- Pluggable routing architecture
  - HAProxy Router
  - F5 Router
- Multiple-routers with traffic sharding
- Router supported protocols
  - HTTP/HTTPS
  - WebSockets
  - TLS with SNI
- Non-standard ports via cloud load-balancers, external IP, and NodePort

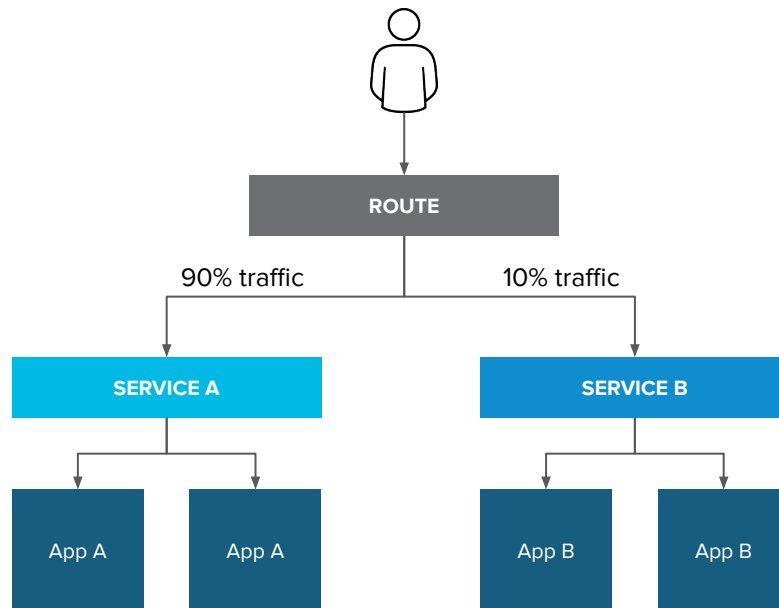


## Routes vs Ingress

Feature	Ingress	Route
Standard Kubernetes object	X	
External access to services	X	X
Persistent (sticky) sessions	X	X
Load-balancing strategies (e.g. round robin)	X	X
Rate-limit and throttling	X	X
IP whitelisting	X	X
TLS edge termination	X	X
TLS re-encryption	X	X
TLS passthrough	X	X
Multiple weighted backends (split traffic)		X
Generated pattern-based hostnames		X
Wildcard domains		X

## Router-based deployment methodologies

Split Traffic Between  
Multiple Services For A/B  
Testing, Blue/Green and  
Canary Deployments

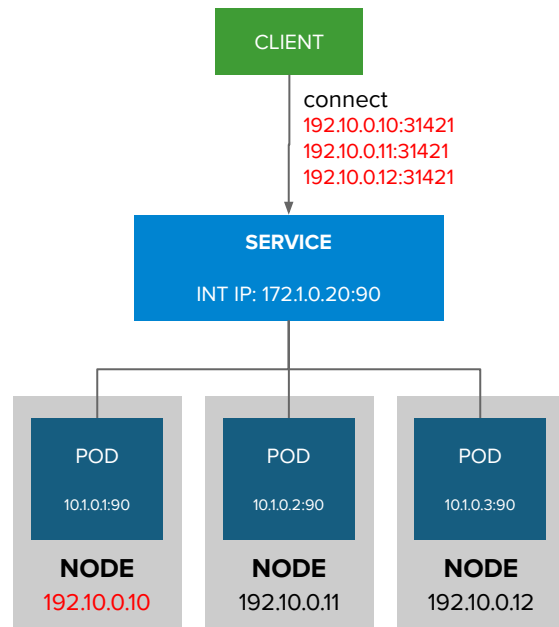


# Alternative methods for ingress

Different ways that traffic can enter the cluster without the router

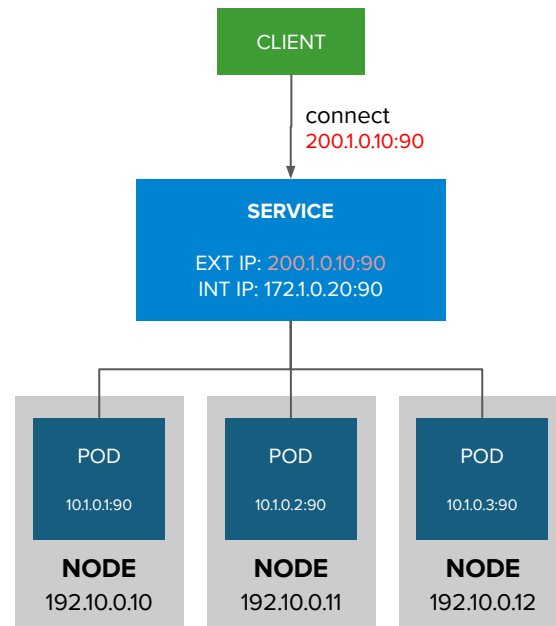
# Entering the cluster on a random port with service nodeports

- NodePort binds a service to a unique port on all the nodes
- Traffic received on any node redirects to a node with the running service
- Ports in 30K-60K range which usually differs from the service
- Firewall rules must allow traffic to all nodes on the specific port



# External traffic to a service on any port with external IP

- Access a service with an external IP on any TCP/UDP port, such as
  - Databases
  - Message Brokers
- Automatic IP allocation from a predefined pool using Ingress IP Self-Service
- IP failover pods provide high availability for the IP pool (fully supported in 4.8)



# Cluster DNS

An automated system  
for providing hostname  
resolution within  
kubernetes



# CoreDNS

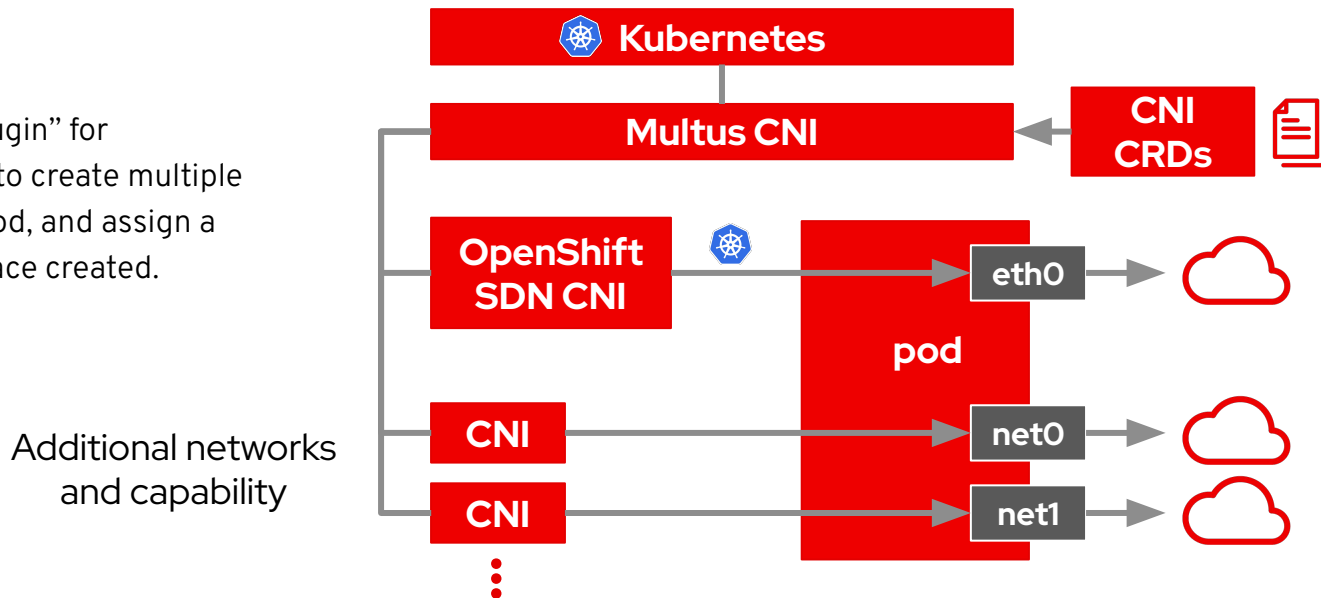
- Built-in internal DNS to reach services by a (fully qualified) hostname
- Split DNS is used with CoreDNS
  - CoreDNS answers DNS queries for internal/cluster services
  - Other defined “upstream” name servers serve the rest of the queries

# Multus

A CNI plugin that  
provides multiple  
network interfaces for  
pods

# Multinetwork with Multus

The Multus CNI “meta plugin” for Kubernetes enables one to create multiple network interfaces per pod, and assign a CNI plugin to each interface created.

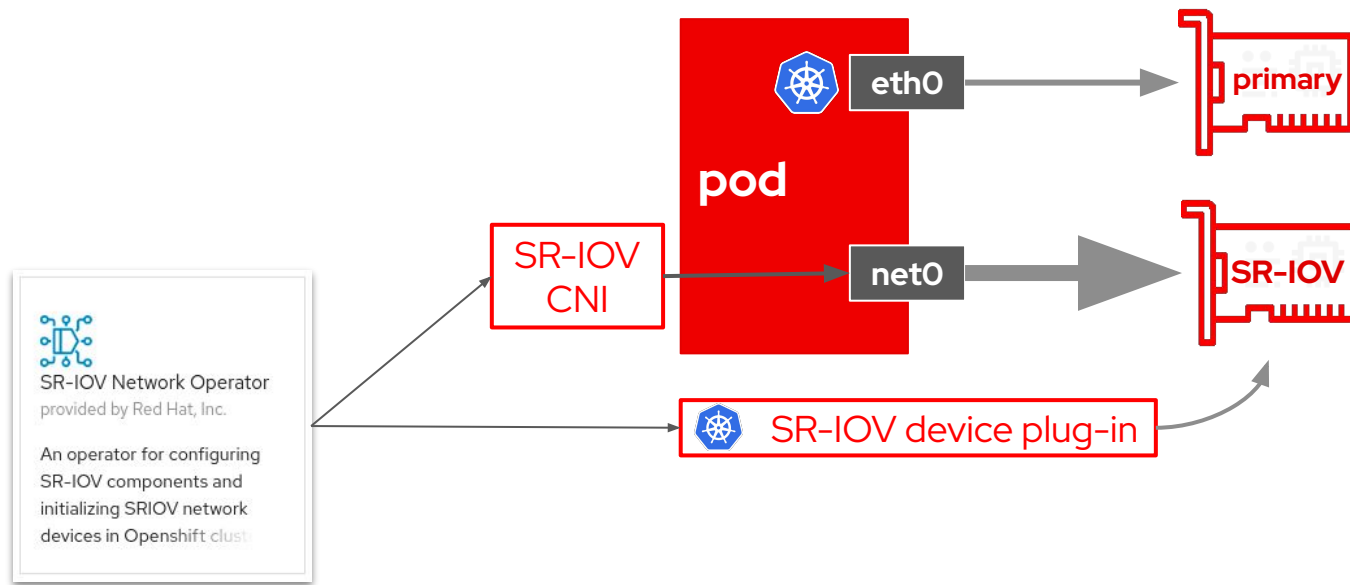


## Additional OpenShift-Supported Secondary CNI Plug-Ins

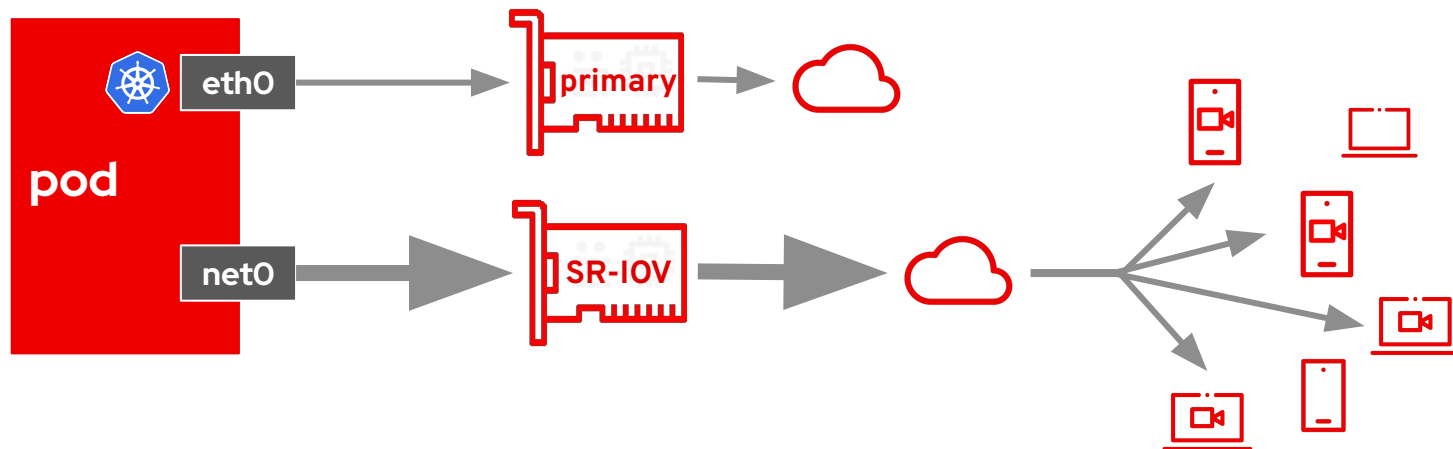
OpenShift 4.x Tested Integrations: [Network Components and Plugins](#)

- host device
- IPAM(dhcp)
- MACVLAN
- IPVLAN
- Bridge with VLAN
- Static IPAM
- DHCP IPAM
- Route Override
- whereabouts
- SR-IOV
- ...

# SR-IOV



# High-performance multicast



# OpenShift Monitoring

An integrated cluster  
monitoring and alerting  
stack

# OpenShift Cluster Monitoring



**Metrics collection and storage**  
via Prometheus, an  
open-source monitoring system  
time series database.

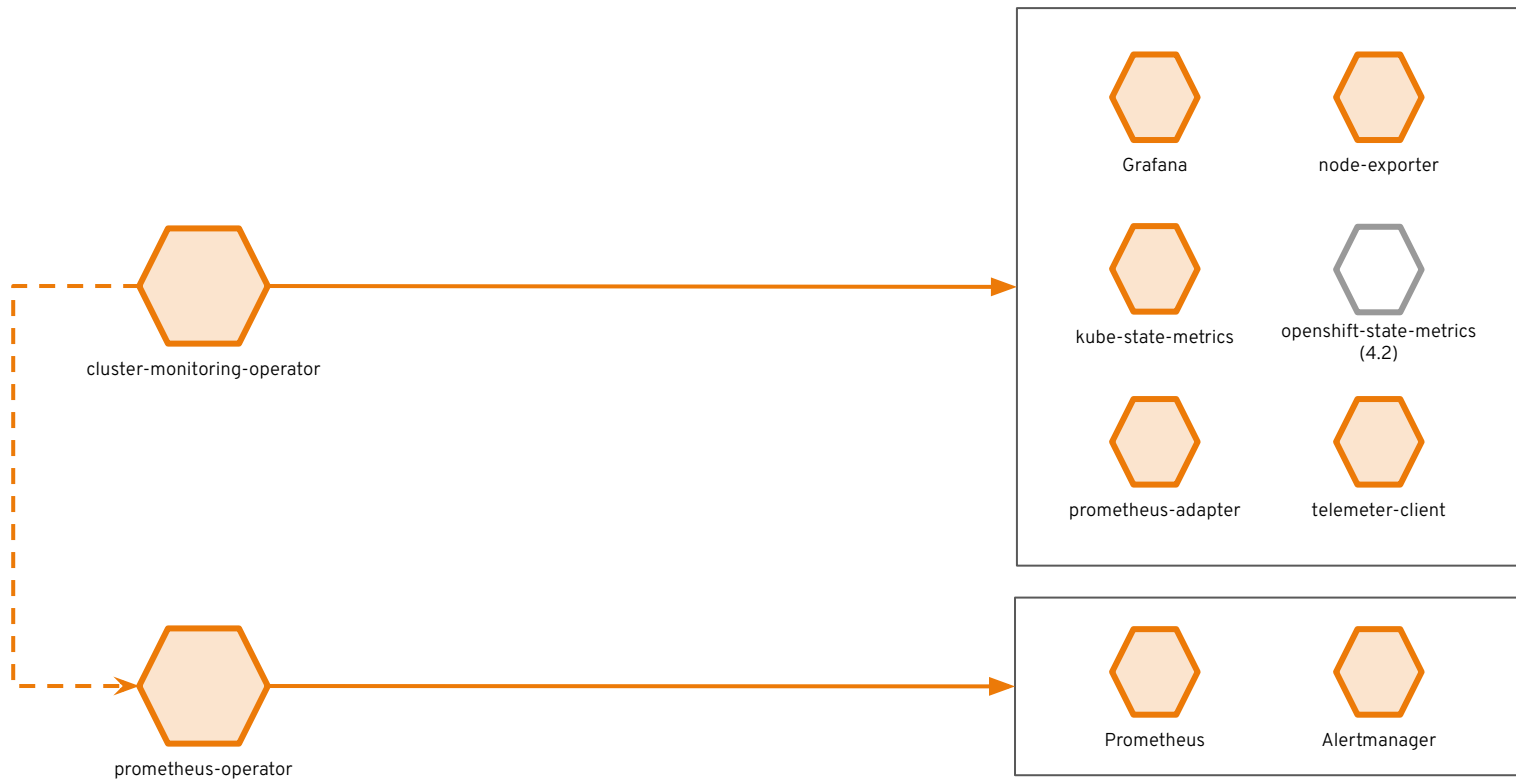


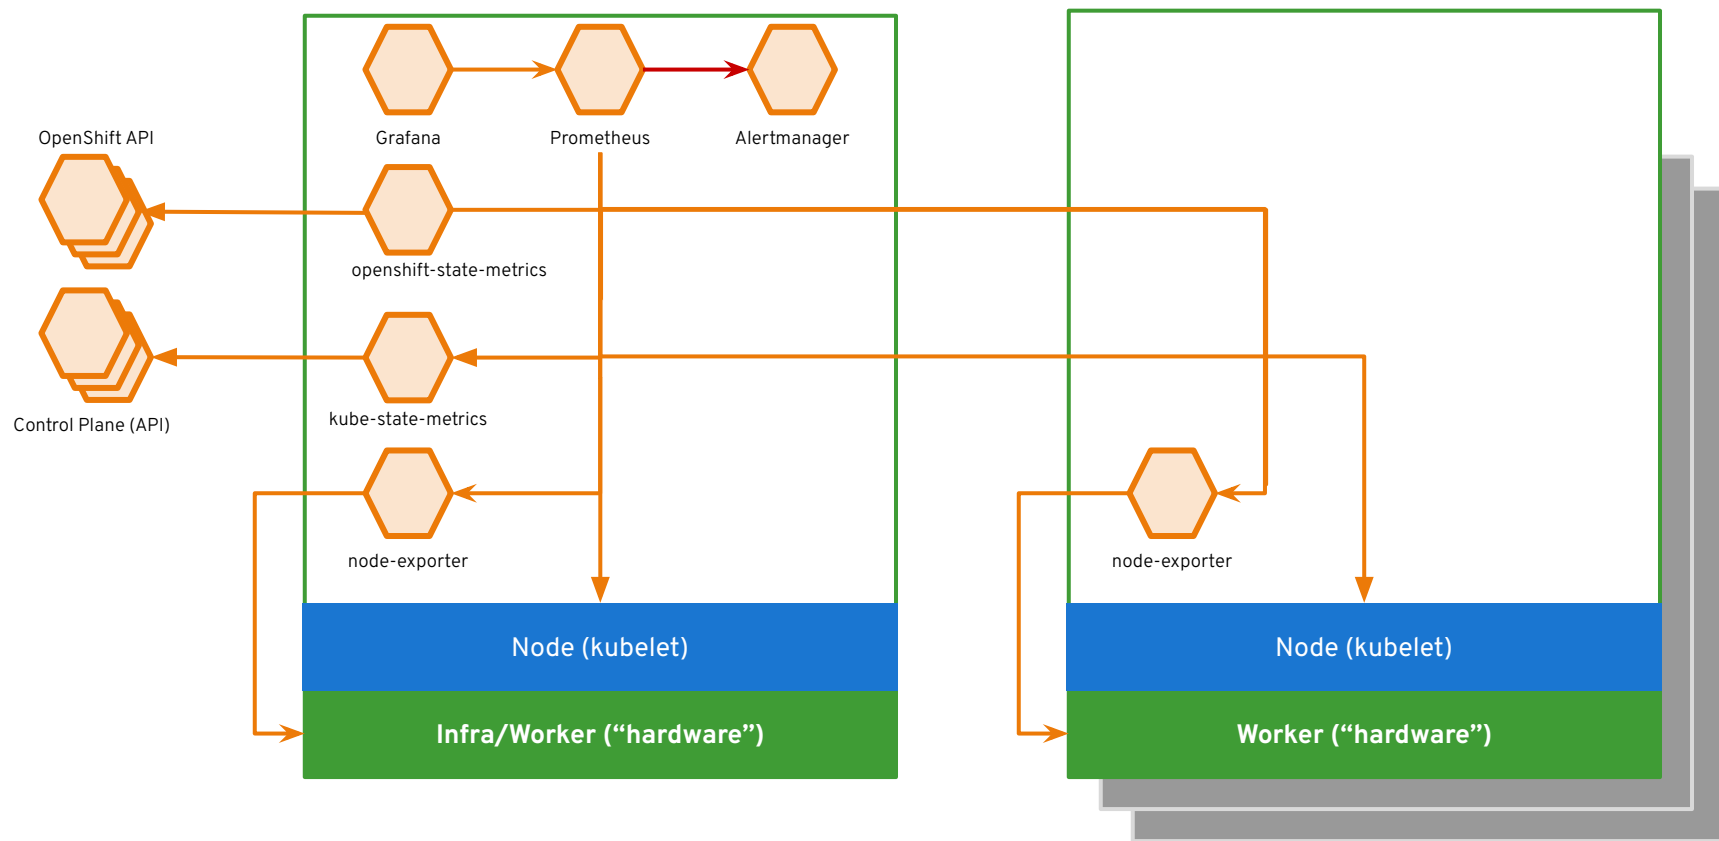
**Alerting/notification** via  
Prometheus' Alertmanager, an  
open-source tool that handles  
alerts send by Prometheus.



**Metrics visualization** via  
Grafana, the leading metrics  
visualization technology.







# OpenShift Logging

An integrated solution  
for exploring and  
corroborating  
application logs

# Observability via log exploration and corroboration with EFK

## Components

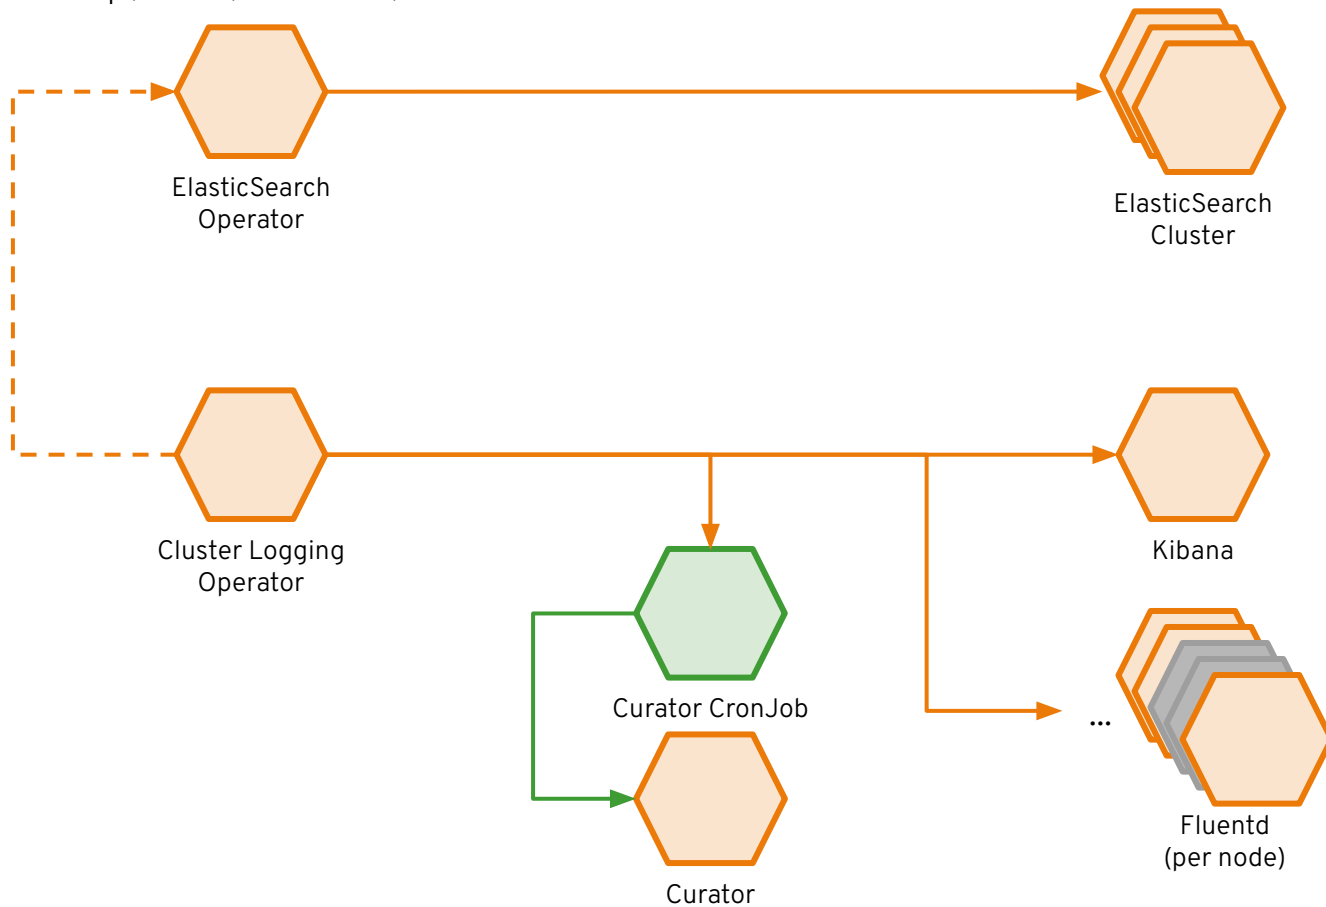
- **Elasticsearch:** a search and analytics engine to store logs
- **Fluentd:** gathers logs and sends to Elasticsearch.
- **Kibana:** A web UI for Elasticsearch.

## Access control

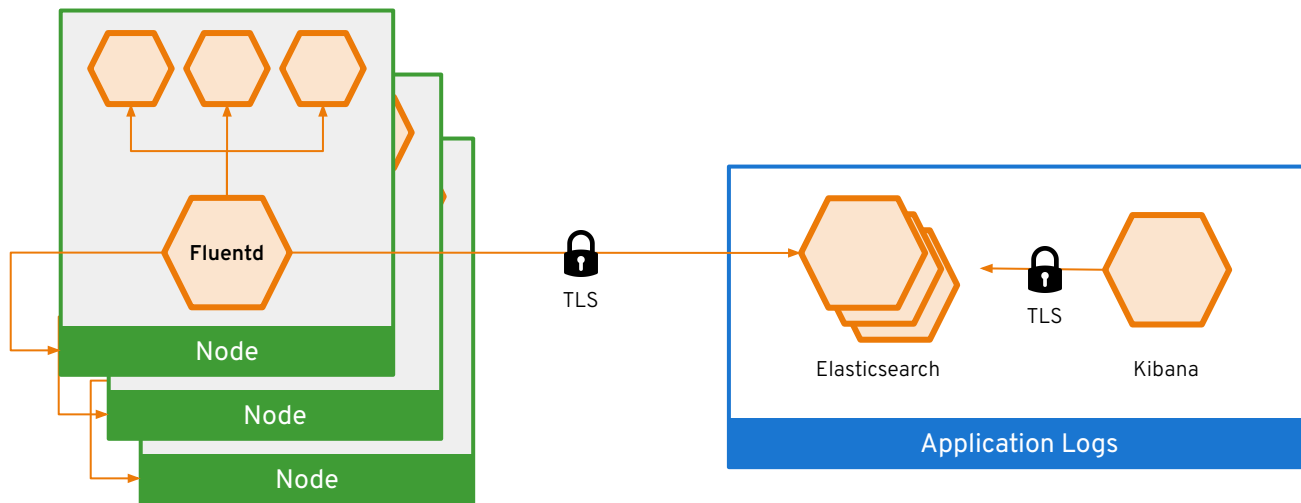
- Cluster administrators can view all logs
- Users can only view logs for their projects

## Ability to forward logs elsewhere

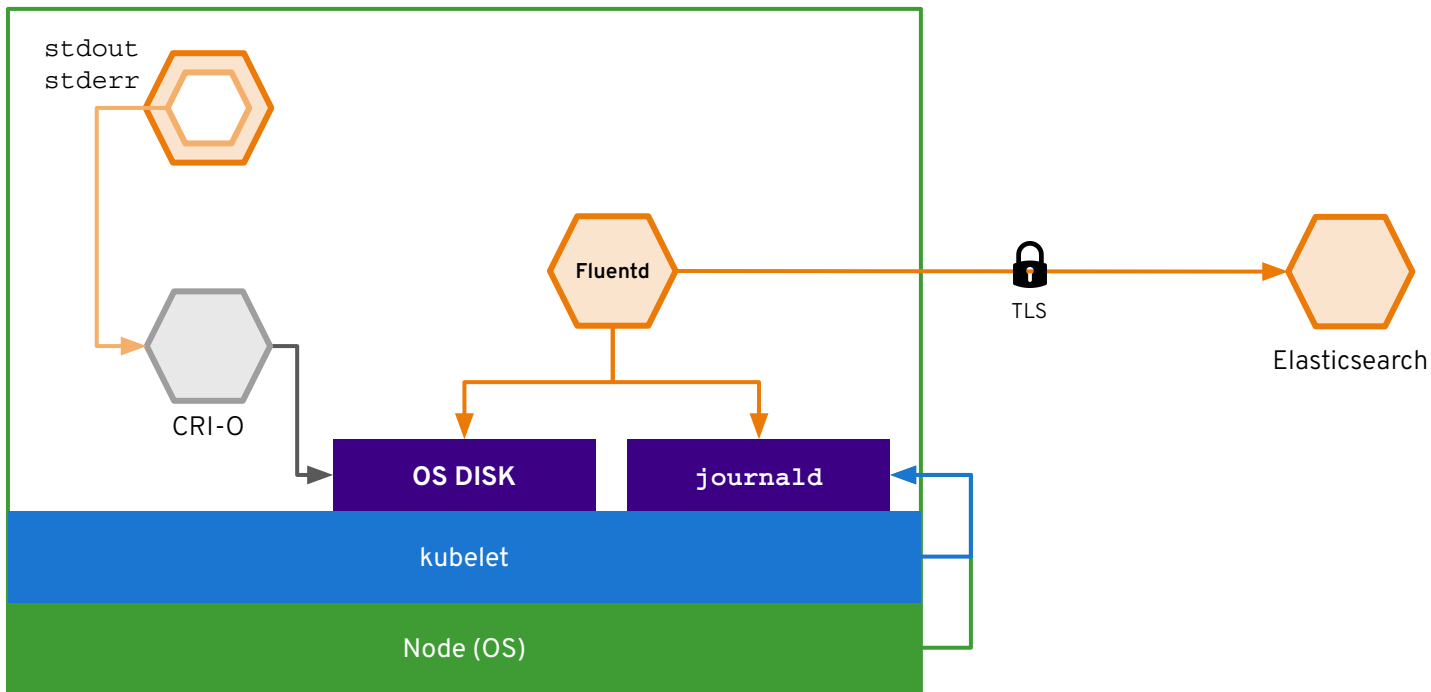
- External elasticsearch, Splunk, etc



## Log data flow in OpenShift



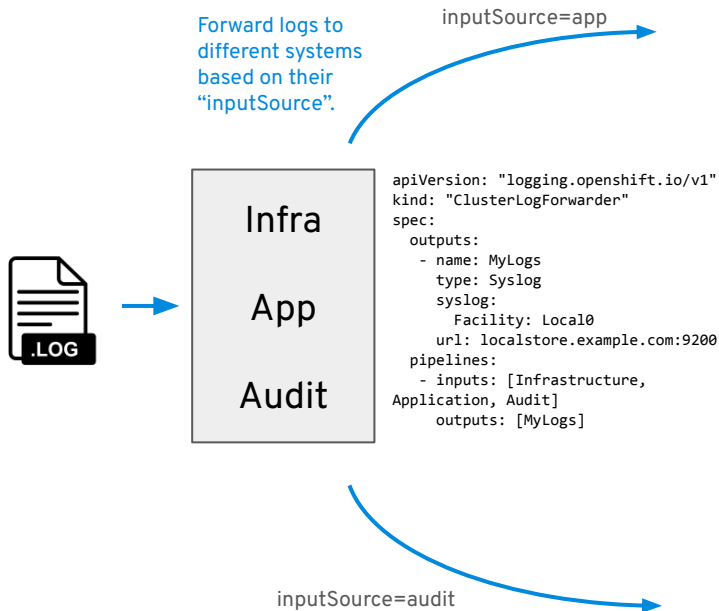
## Log data flow in OpenShift



# New log forwarding API (since 4.6)

**Abstract Fluentd configuration by introducing new log forwarding API to improve support and experience for customers.**

- Introducing a new, cluster-wide *ClusterLogForwarder* CRD (API) that replaces needs to configure log forwarding via Fluentd ConfigMap.
- The API helps to reduce probability to misconfigure Fluentd and helps bringing in more stability into the Logging stack.
- Features include: Audit log collection and forwarding, Kafka support, namespace- and source-based routing, tagging, as well as improvements to the existing log forwarding features (e.g. syslog RFC5424 support).





# Secure Log Forwarding to 3rd party

```
apiVersion: "logging.openshift.io/v1"
kind: "ClusterLogForwarder"
spec:
  outputs:
    - name: MyLogs
      type: Syslog
      syslog:
        Facility: Local0
        url: localstore.example.com:9200
  pipelines:
    - inputs: [Infrastructure,
      Application, Audit]
      outputs: [MyLogs]
```

"ClusterLogForwarder"  
Custom Resource

Cluster Logging  
Operator

watches

creates

elasticsearch

fluentd

kafka

SYS  
LOG

External Logging system

Node

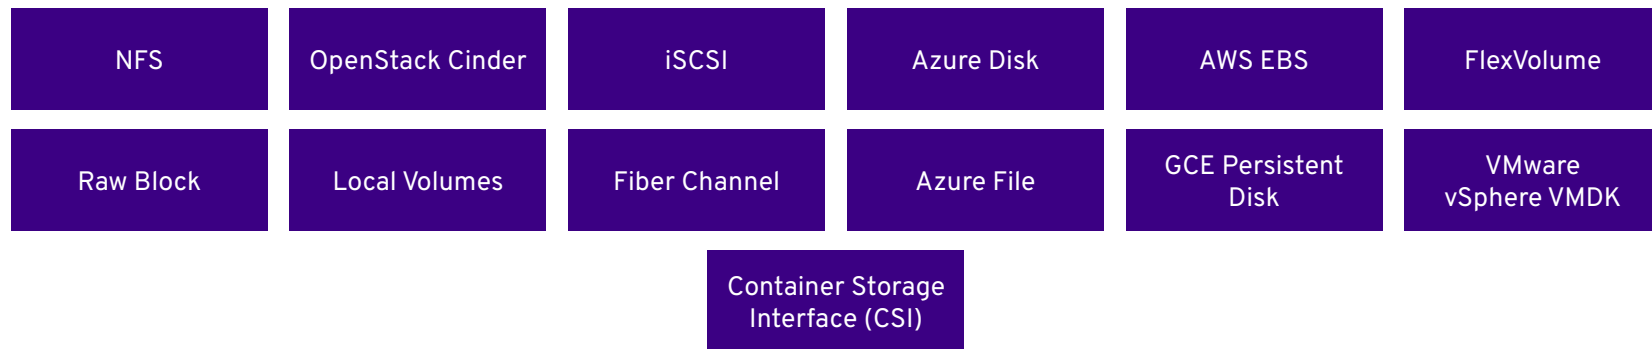
Fluentd  
daemonset

Fluentd  
forwarder

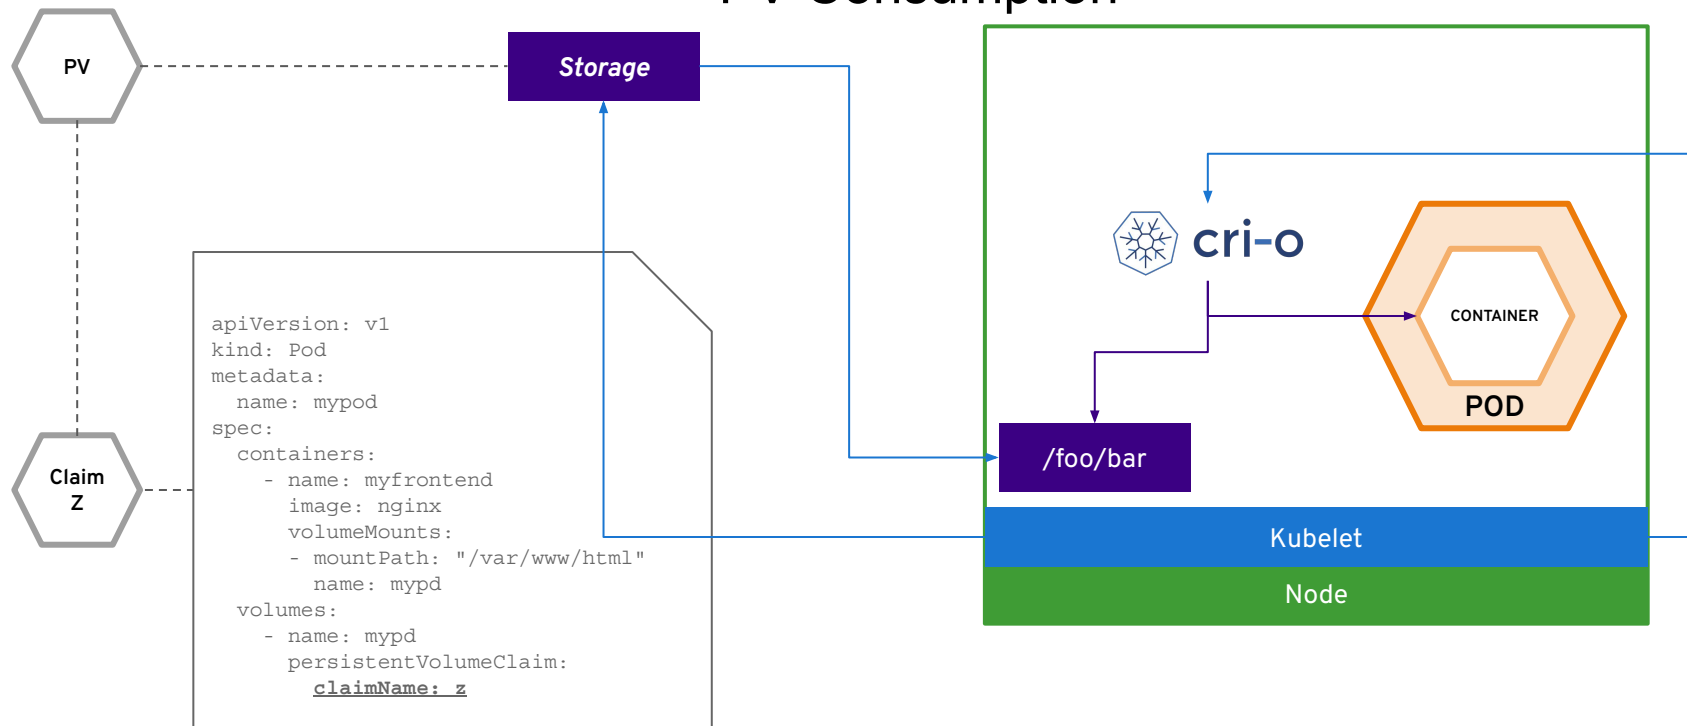
# Persistent Storage

Connecting real-world  
storage to your  
containers to enable  
stateful applications

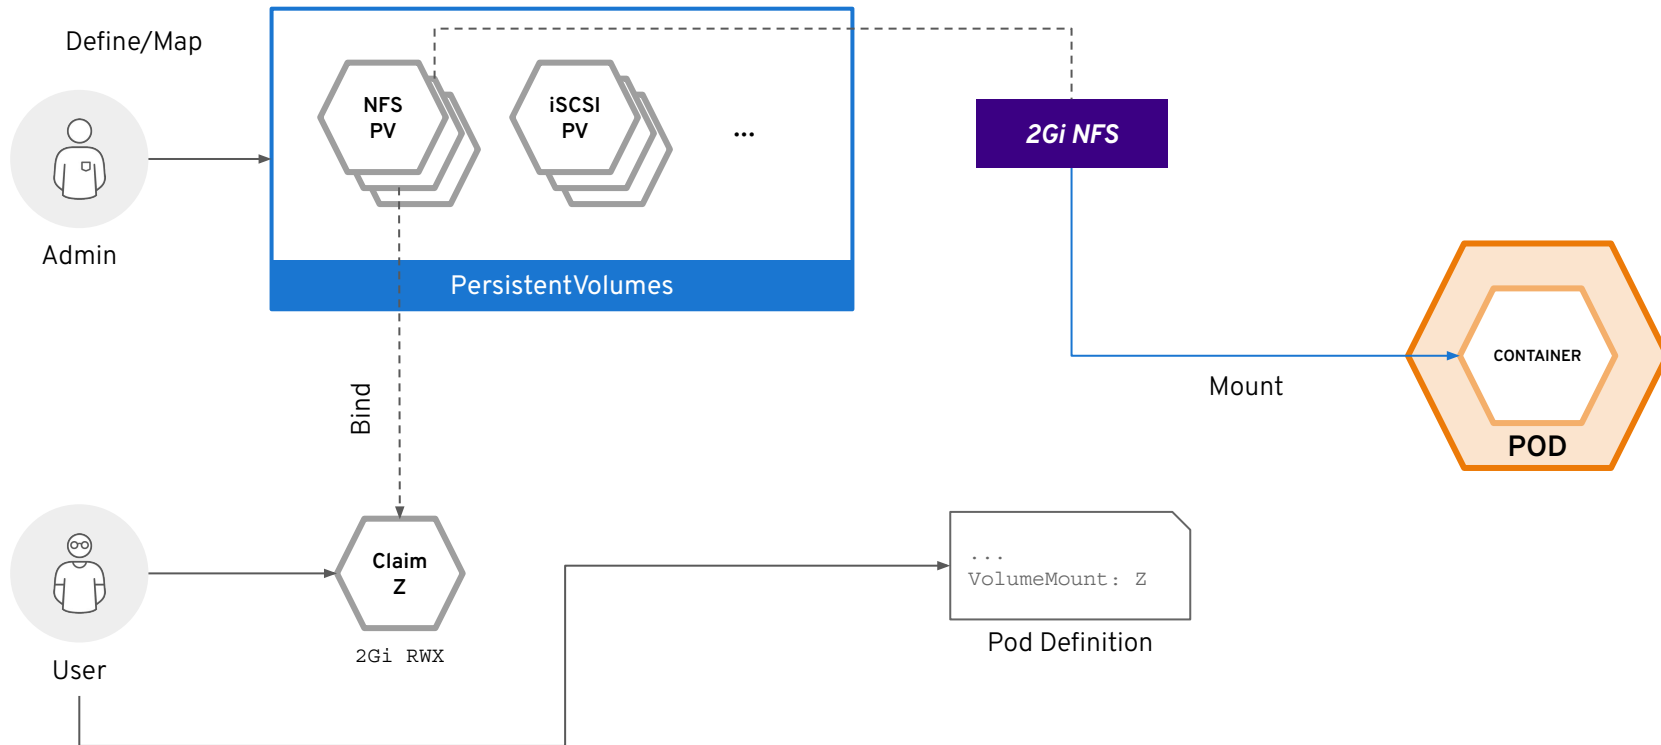
## A broad spectrum of static and dynamic storage endpoints



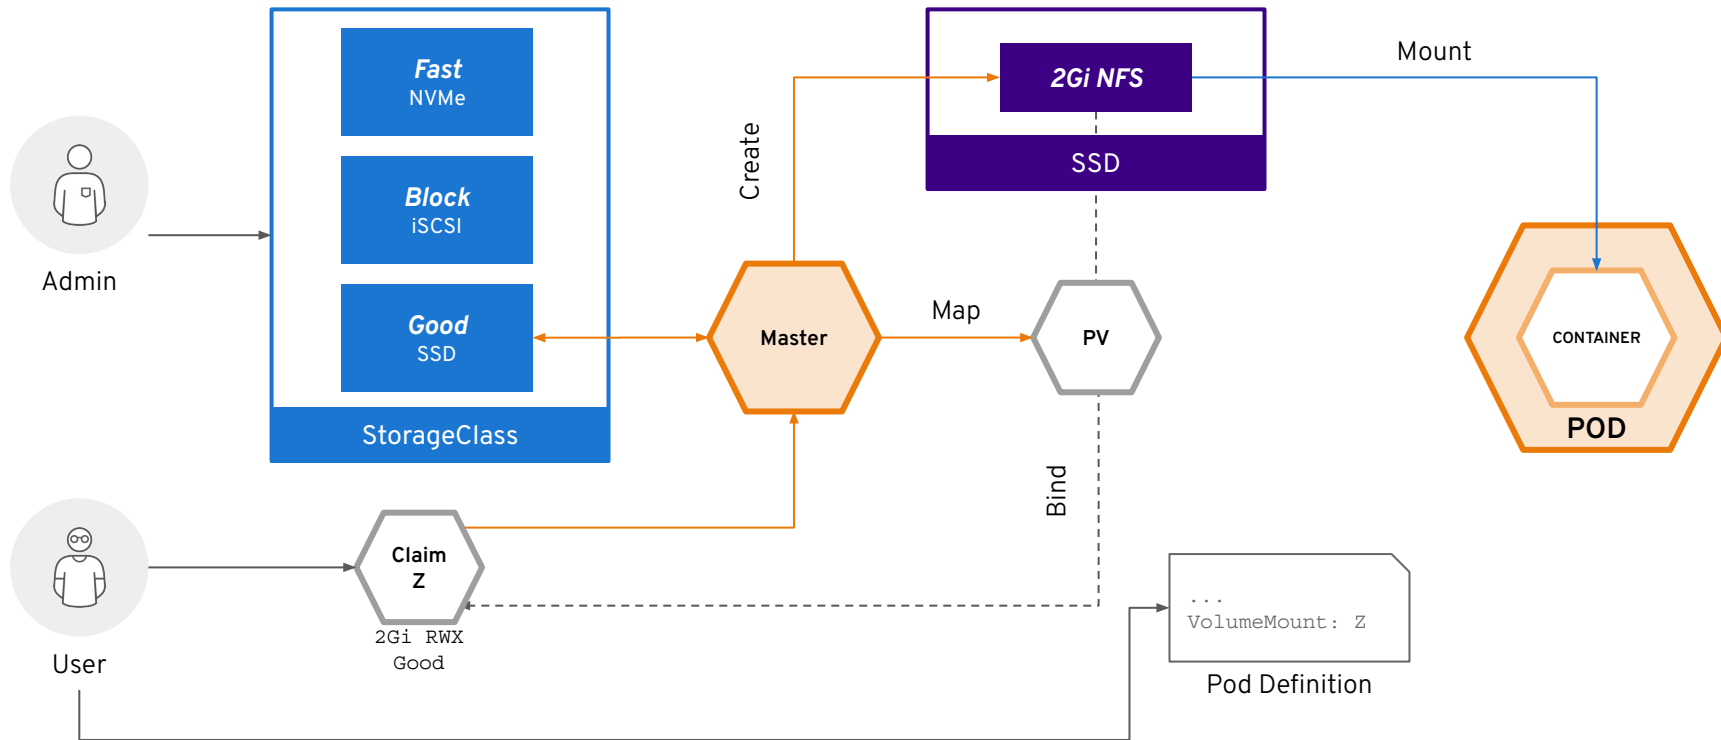
## PV Consumption



## Static Storage Provisioning

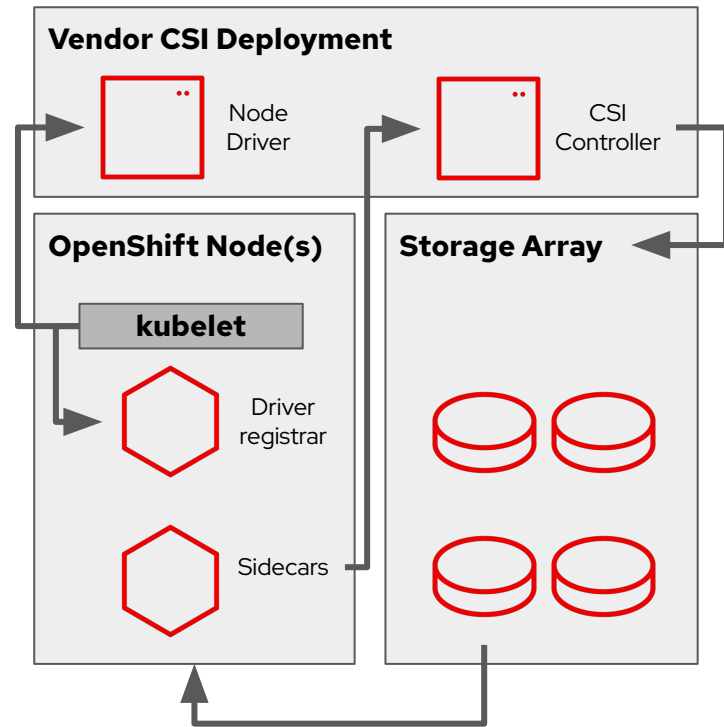


# Dynamic Storage Provisioning



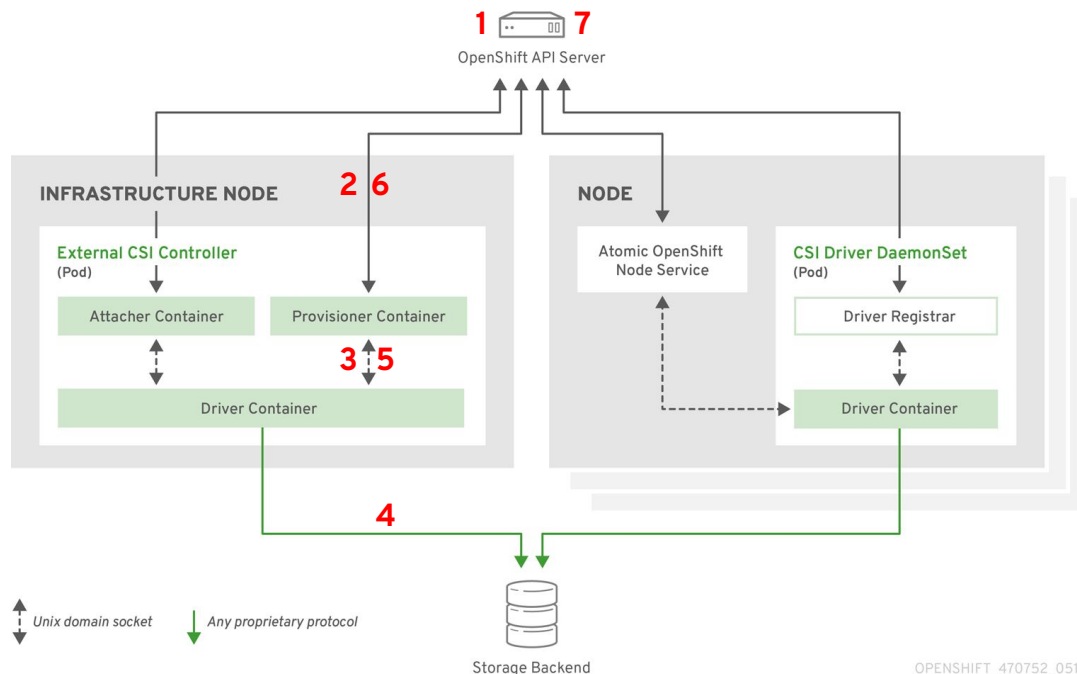
## CSI Driver Paradigm

- CSI drivers and logic are provided by storage vendors
  - Each implementation may be different based on the vendor
- Controller logic is deployed to the OpenShift cluster as an Operator, deployment, or even a standalone Pod(s)
  - Responsible for interfacing with storage device to create and manage volumes, snapshots, clones, etc.
  - Respond to events (create, delete PVC) for assigned StorageClass(es)
  - Sidecars assist with hooks for additional functionality - snapshots, resizing, etc.
- Each node hosts, via a DaemonSet, one or more CSI node plugin Pods for the driver
  - Kubelet requests the node plugin to mount/unmount volumes, format block devices if needed, etc.



# CSI Dynamic Provisioning

1. User creates a PVC
2. The external provisioner gets an event that a new PVC was created
3. The external provisioner initiates CreateVolume call to the CSI driver
4. The CSI driver talks to storage backend and creates a volume
5. The CSI driver returns a volume to the external provisioner
6. The external provisioner creates PV on API server
7. Kubernetes PV controller finishes the binding (PVC is Bound)

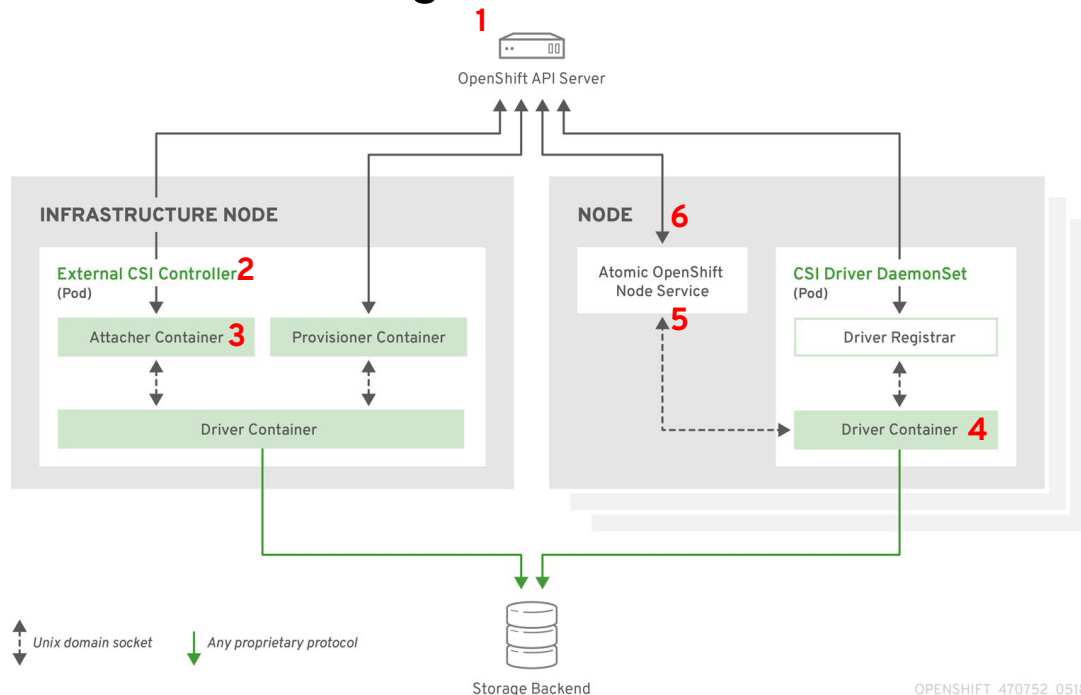


OPENSIFT\_470752\_0518



# CSI Volume Mounting

1. User instantiates a Pod with a PVC
2. The CSI controller is notified of a volume publish event via the attacher sidecar
3. The CSI controller takes any actions on the storage device to make the volume mountable, e.g. NFS export rules
4. The node driver stages the volume, taking action to prepare the volume to be used, e.g. formatting a non-raw block device
5. The node driver mounts the volume at the location requested by Kubelet
6. The volume is attached to the container, by Kubelet, as defined



OPENSHIFT\_470752\_0518

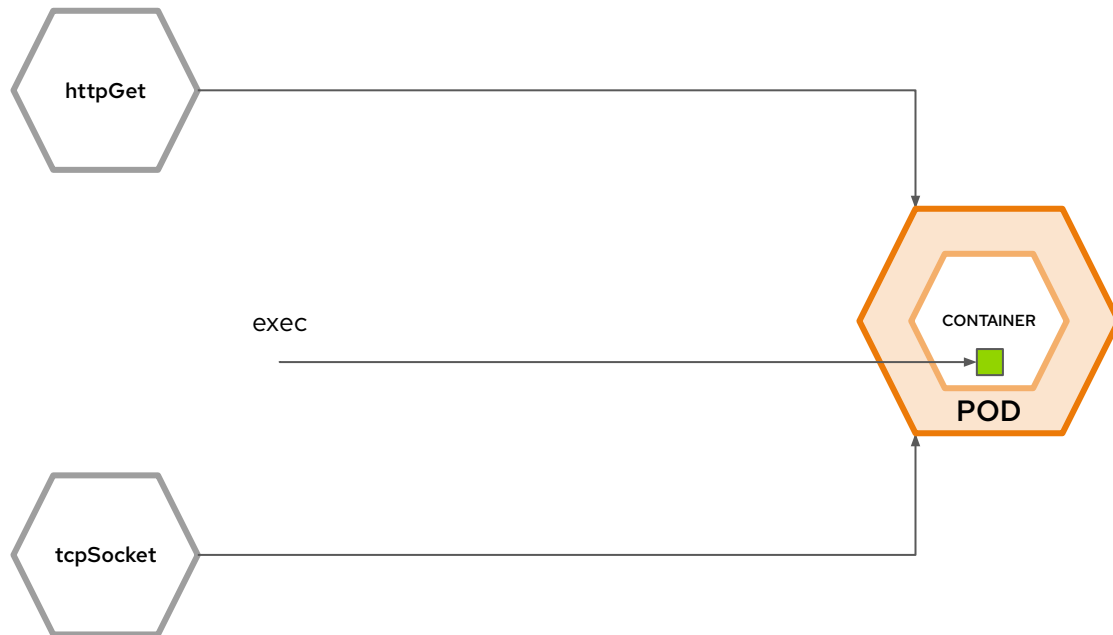


# Developer Experience

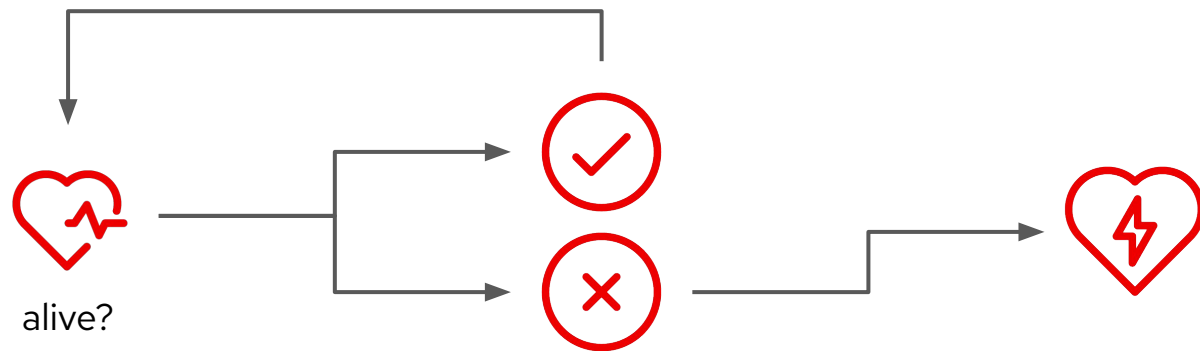
# Application Probes

Improving reliability  
and availability of  
applications via built-in  
probes

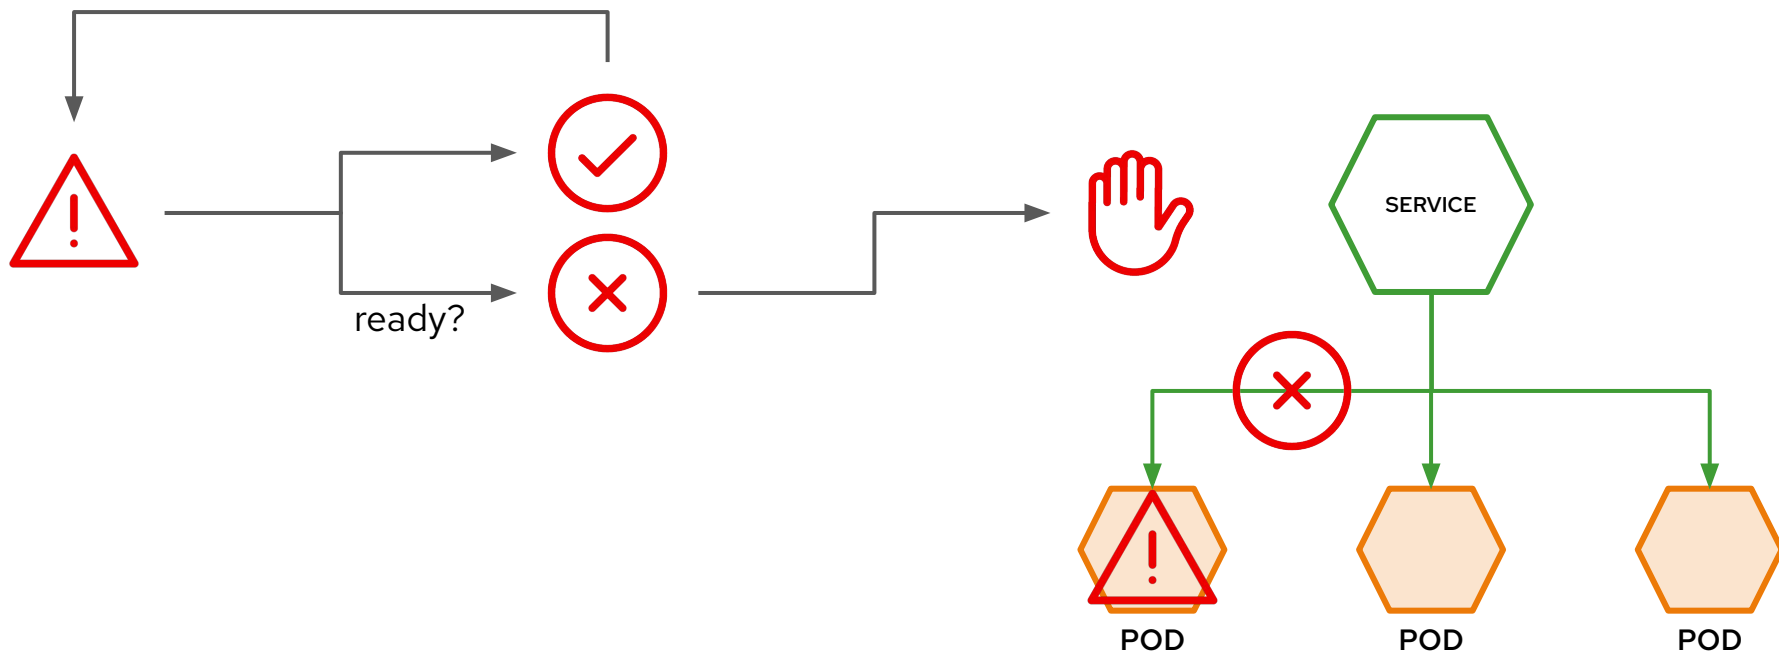
## Three Types: One Goal



## Liveness Probes



## Readiness Probes



## Important settings

`initialDelaySeconds`: How long to wait after the pod is launched to begin checking

`timeoutSeconds`: How long to wait for a successful connection (httpGet, tcpSocket only)

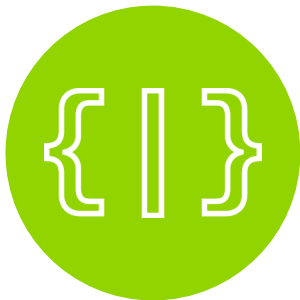
`periodSeconds`: How frequently to recheck

`failureThreshold`: How many consecutive failed checks before the probe is considered failed

# Build and Deploy Container Images

Tools and automation  
that makes developers  
productive quickly





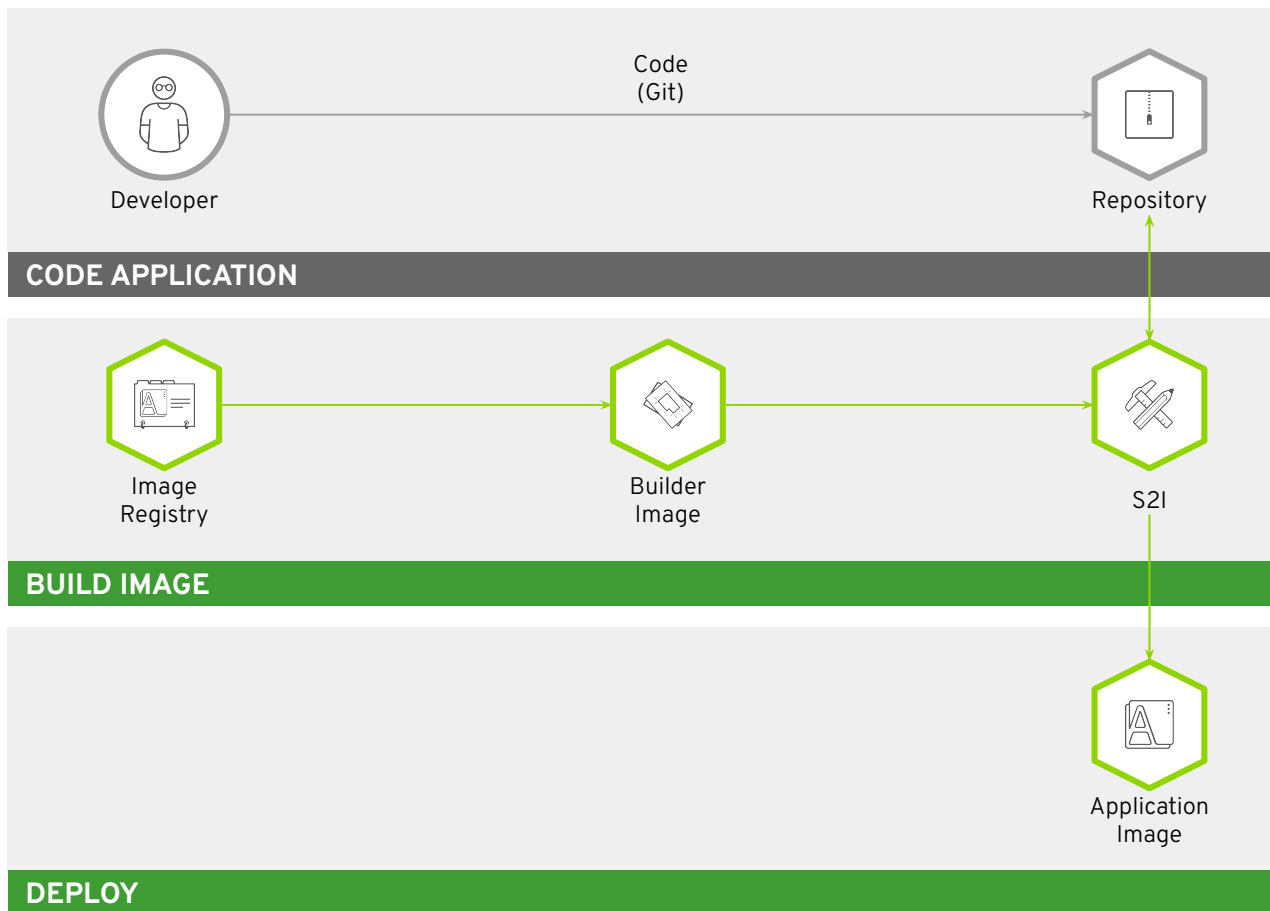
**DEPLOY YOUR  
SOURCE CODE**



**DEPLOY YOUR  
APP BINARY**



**DEPLOY YOUR  
CONTAINER IMAGE**





[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[twitter.com/RedHat](https://twitter.com/RedHat)